Modelling time-varying features of speech: tools and methods

Michele Gubian

Institute of Phonetics and Speech Processing
LMU Munich, Germany

Graz, October 2020

Alignment of rising pitch accents in Spanish

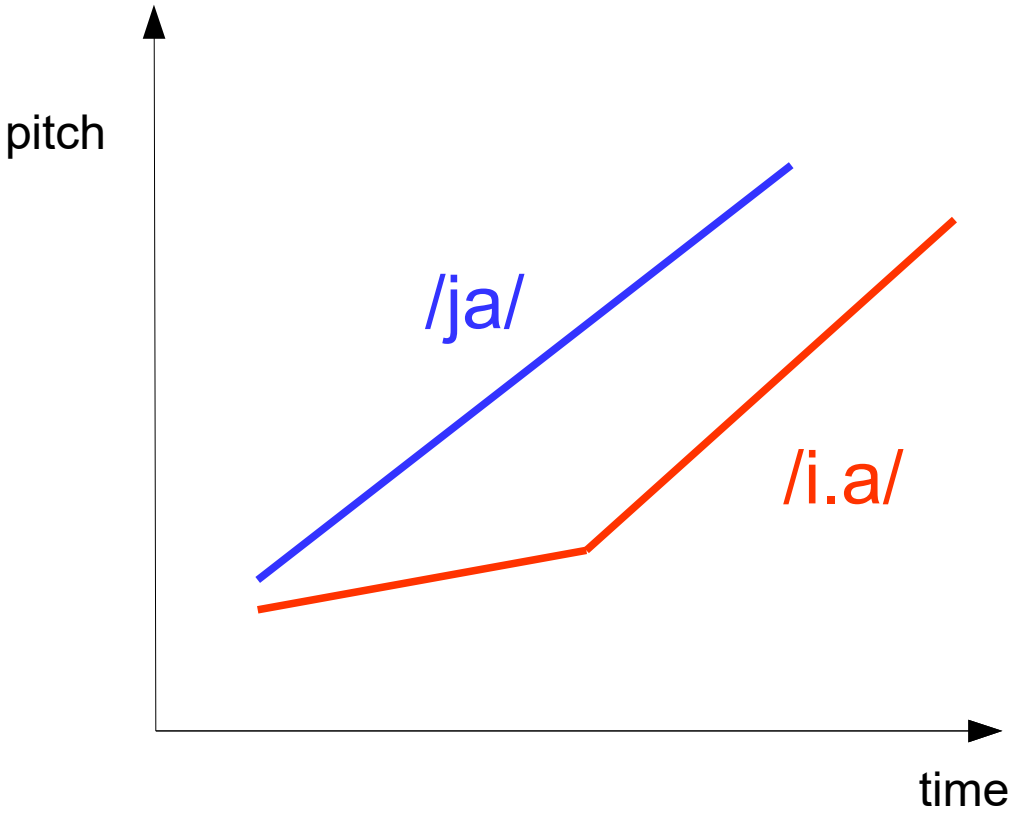

- European Spanish
- **Diphthong**: /ja/

  e.g. *Emili<u>a</u>na*
- **Hiatus** /i.a/

  e.g. *pi<u>a</u>no*
- Rising pitch accent should align to syllabic structure

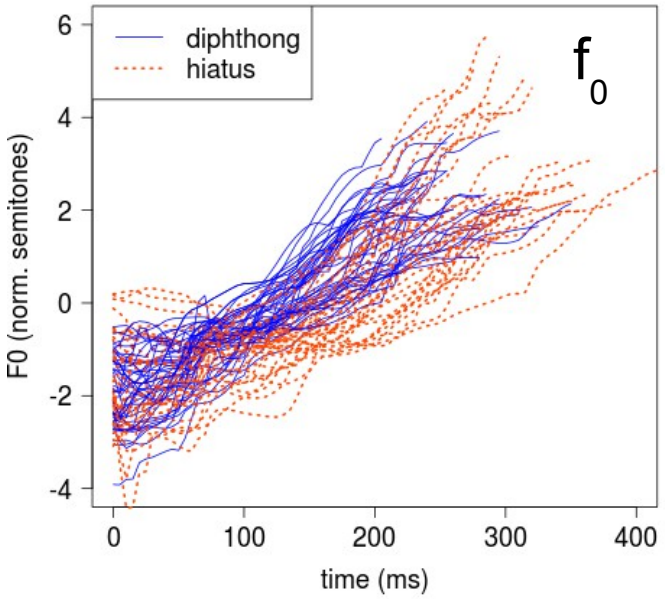

- Read speech
- 9 participants
- 20 Diphthongs +

  20 Hiatuses each

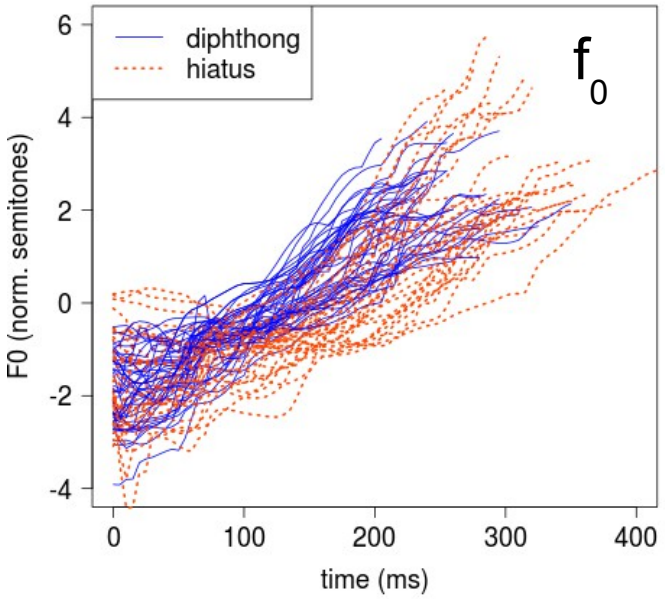

LMER

- Read speech
- 9 participants
- 20 Diphthongs +

  20 Hiatuses each

CURVES

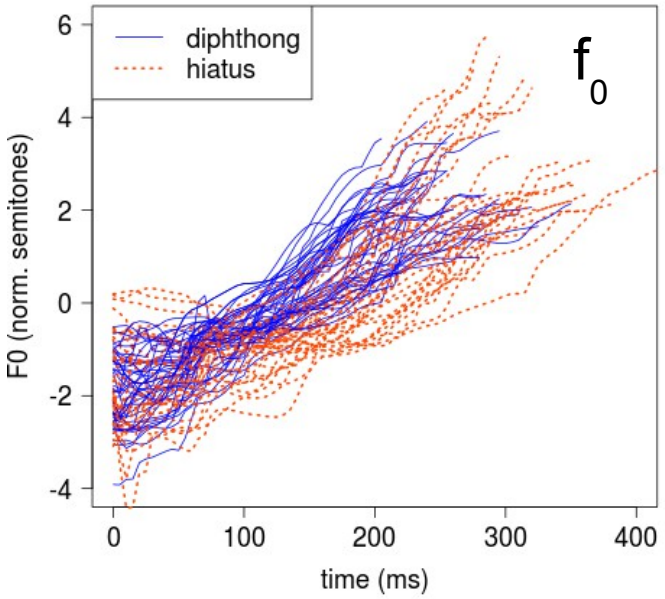

- Read speech
- 9 participants
- 20 Diphthongs +

  20 Hiatuses each

NUMBERS

LMER

CURVES

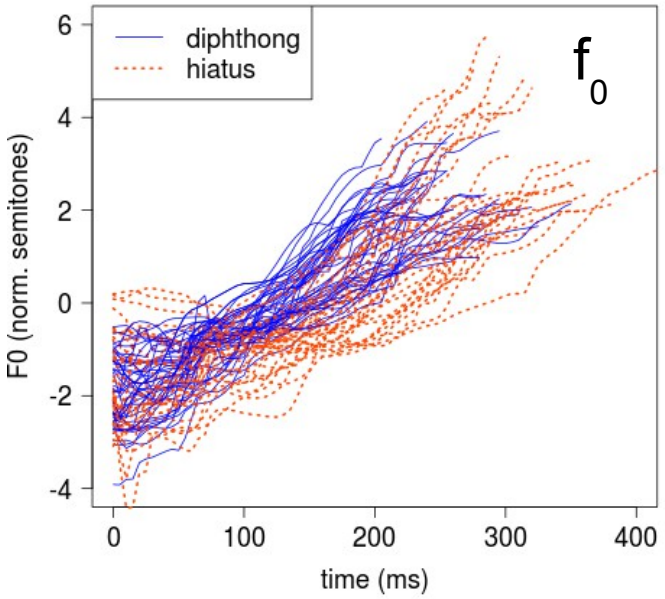

- Read speech
- 9 participants
- 20 Diphthongs +
  20 Hiatuses each

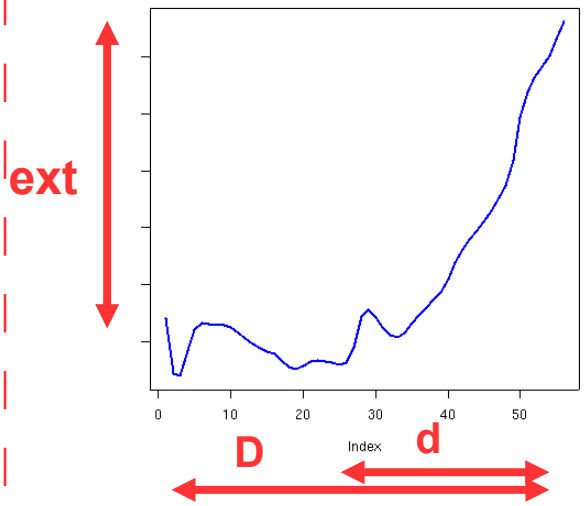

ext (st)	d/D	Cat.
5.3	0.9	D
4.6	0.7	H
...	...	...

NUMBERS

LMER

CURVES

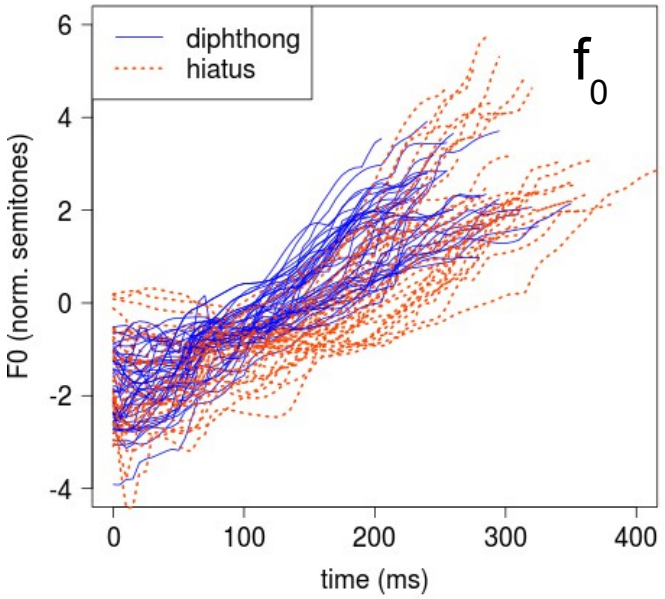

- Read speech
- 9 participants
- 20 Diphthongs +
  20 Hiatuses each

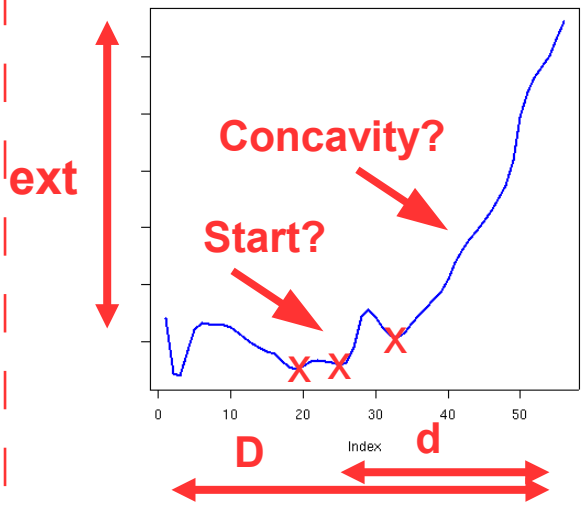

ext (st)	d/D	Cat.
5.3	0.9	D
4.6	0.7	H
…	…	…

NUMBERS

LMER

CURVES

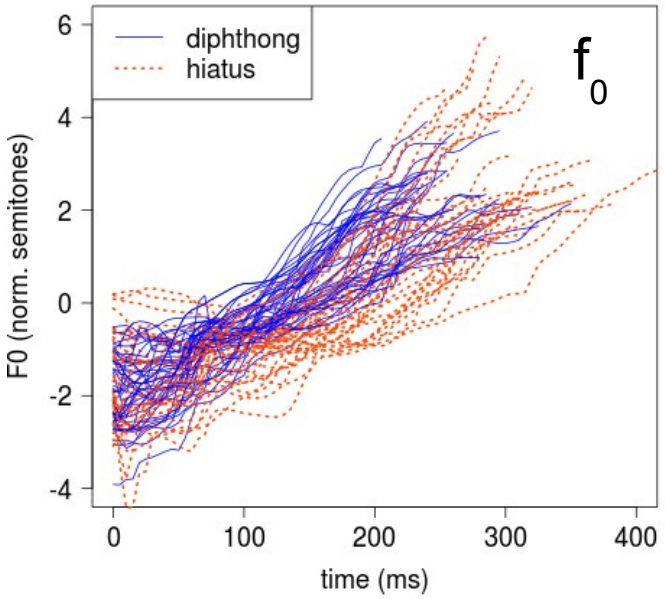

- Read speech
- 9 participants
- 20 Diphthongs +

  20 Hiatuses each

DCT

k0	k1	k2	...
...	...	...	...
...	...	...	...
...	...	...	...

NUMBERS

LMER

DCT limitations

- DCT does not (easily) encode time-localised information, e.g. a small hump
- Typically only k0, k1 and k2 are used, which have a geometric interpretation (mean, slope, curvature)
- Extracting several k's brings up the need of PCA
- In general, not effective to encode long signals

CURVES

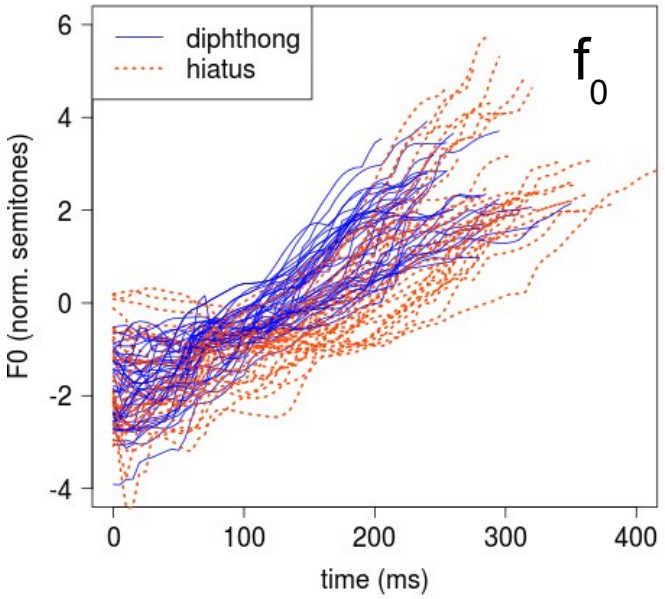

- Read speech
- 9 participants
- 20 Diphthongs +

  20 Hiatuses each

NUMBERS

GAMM

GAMMs

- **PRO**
  - LMER directly on curves
  - Good R packages (e.g. mgcv)
  - Good tutorials (e.g. Wieling, Soskuthy)
- **CON**
  - No easy way to analyse multidimensional signals
  - Computationally heavy
  - LMER directly on curves :D

CURVES

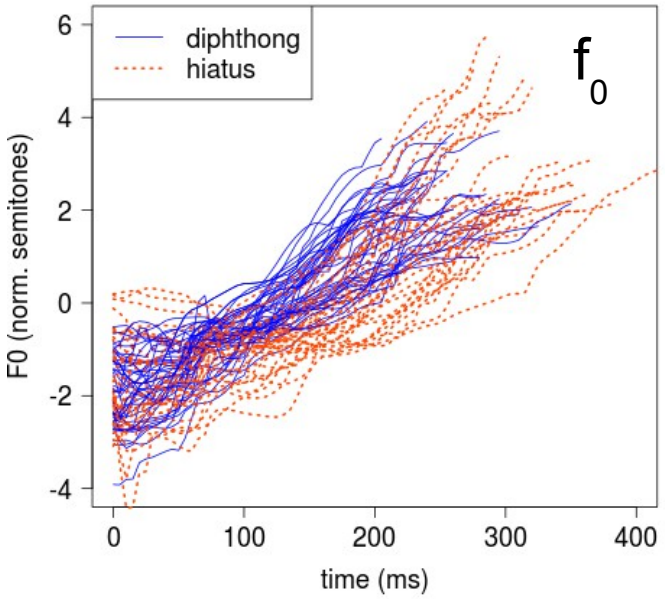

- Read speech
- 9 participants
- 20 Diphthongs +
  20 Hiatuses each

FPCA

s1	s2	s3	...
...	...	...	...
...	...	...	...
...	...	...	...

NUMBERS

LMER

FPCA

- **PRO**
  - Computationally light
  - Interpretable
  - Easy to analyse multidimensional signals
  - Allows you to use LMER
- **CON**
  - Suboptimal with respect to GAMMs
  - Can fail with many categories
  - Can fail when noise is linearly related to variation of interest

Road map

CURVES

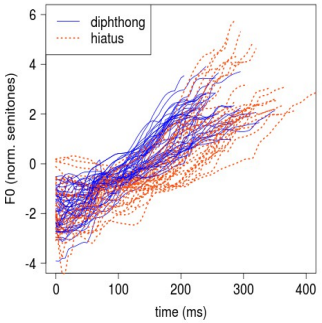

Interpolate using a function basis

- Data driven

Dimensionality reduction tool

- Few parameters
- Interpretable

NUMBERS

LMER

Road map

CURVES

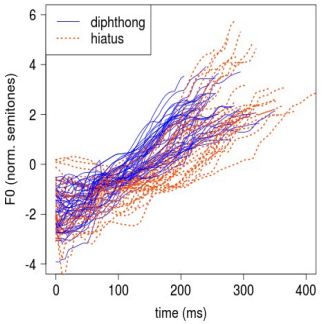

Interpolate using a function basis

- Data driven

Dimensionality reduction tool

- Few parameters
- Interpretable

NUMBERS

LMER

Interpolation with B-splines

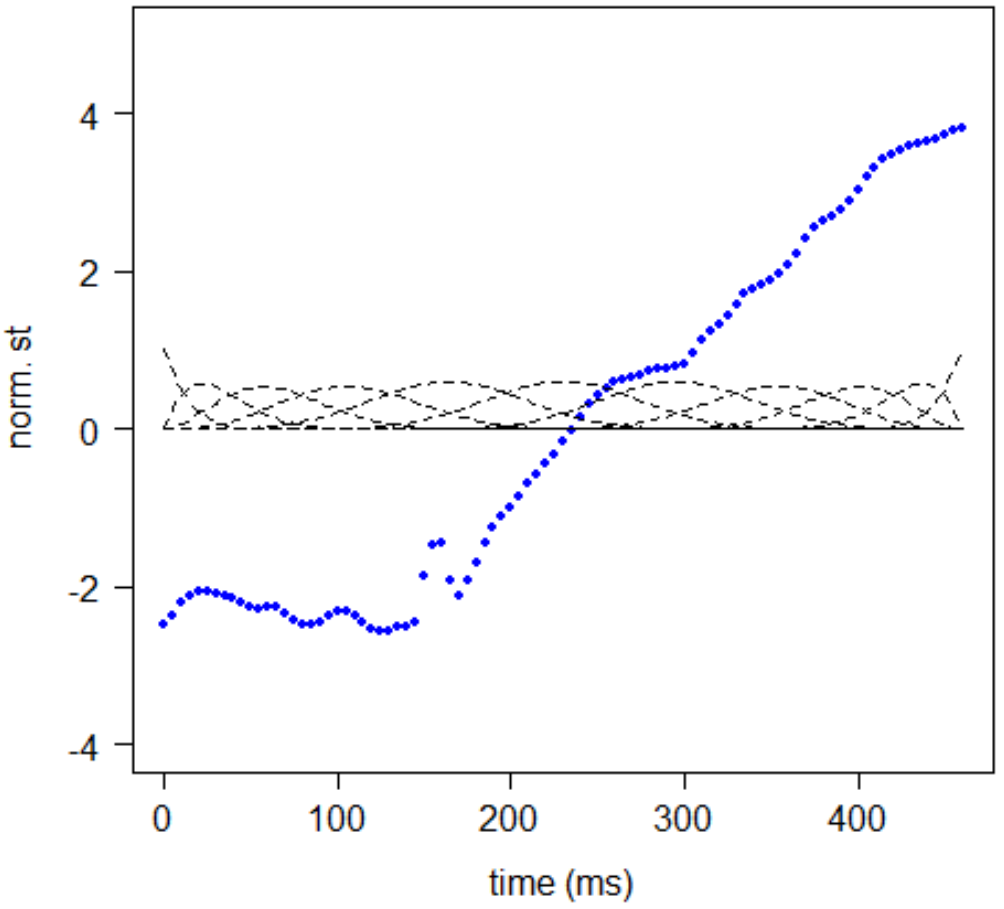

Interpolation with B-splines

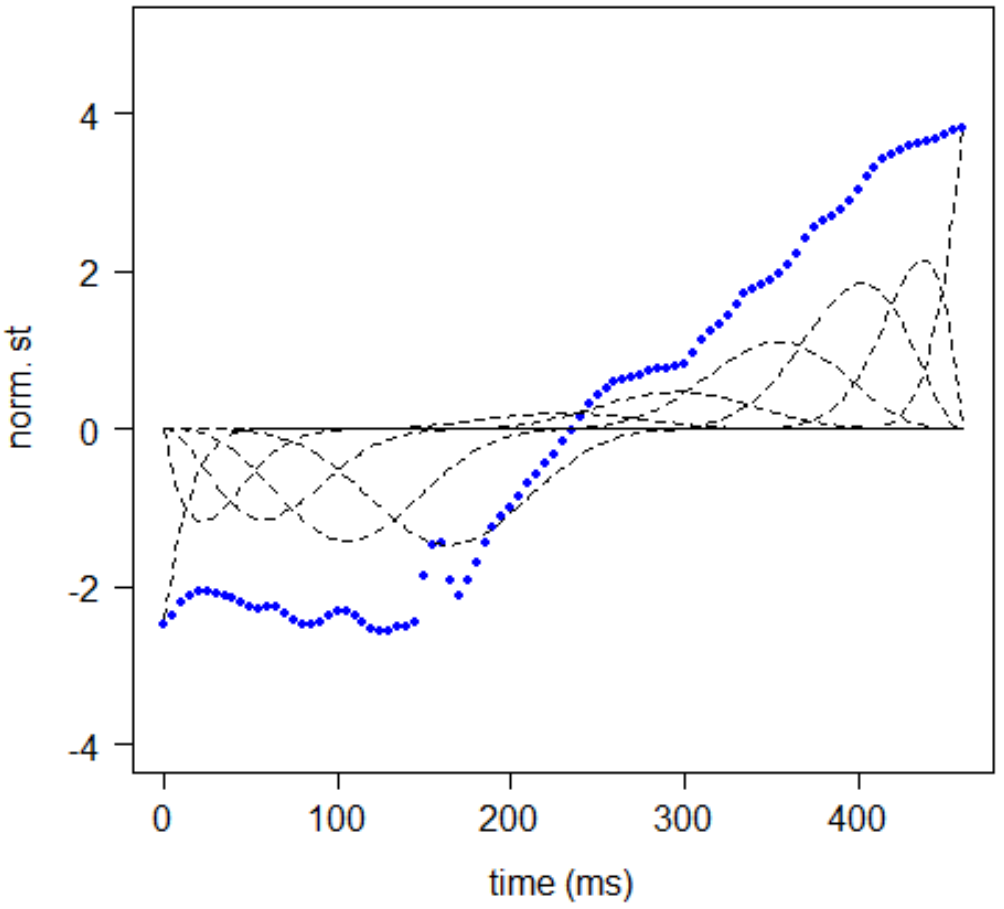

Interpolation with B-splines

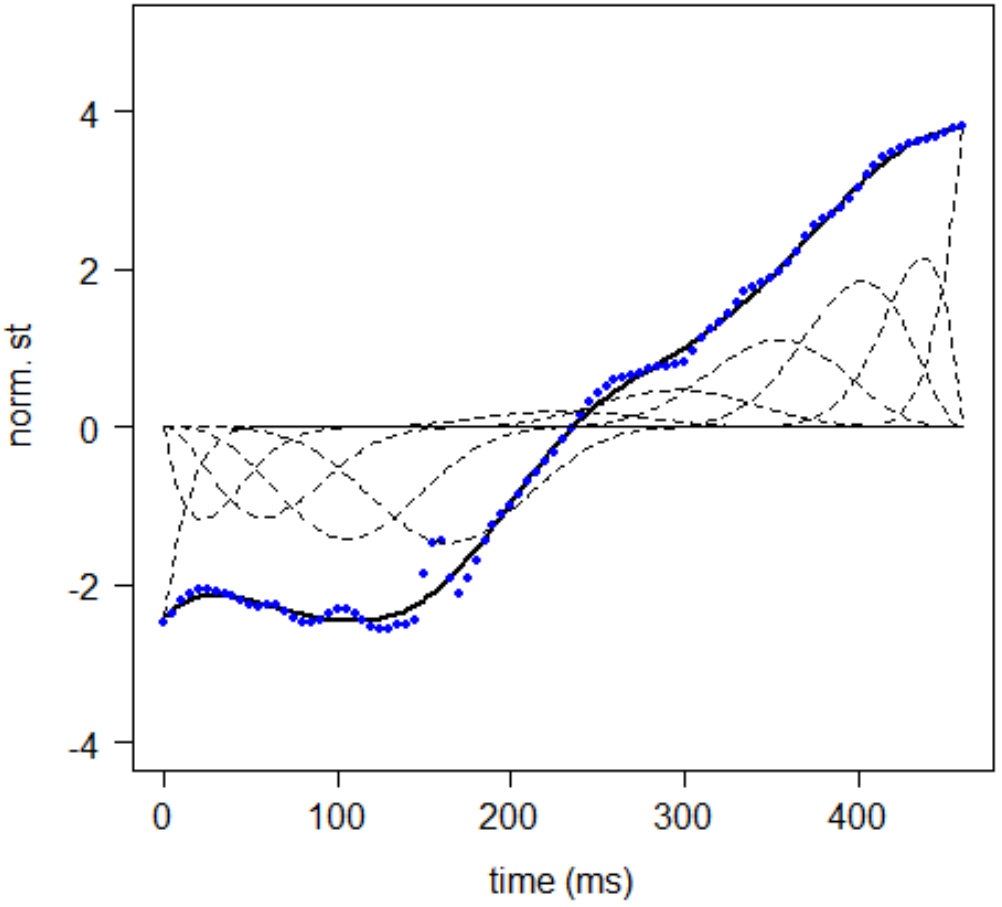

Different durations

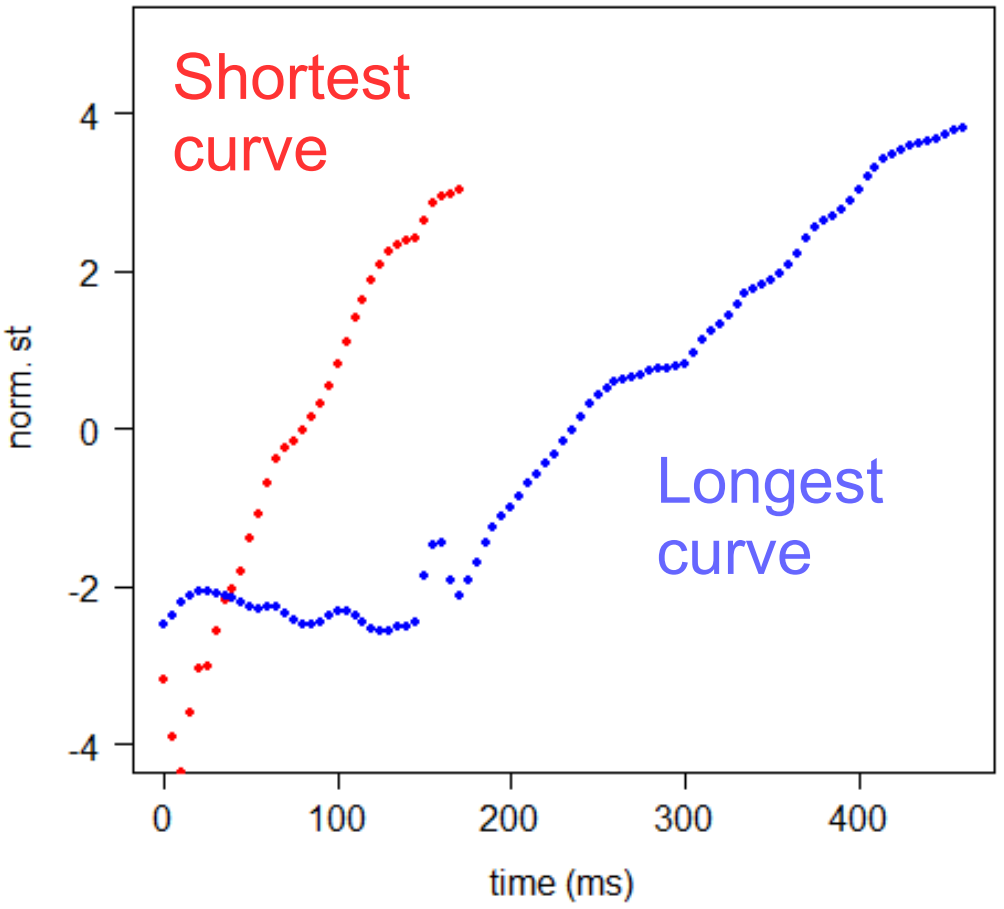

Linear time normalisation

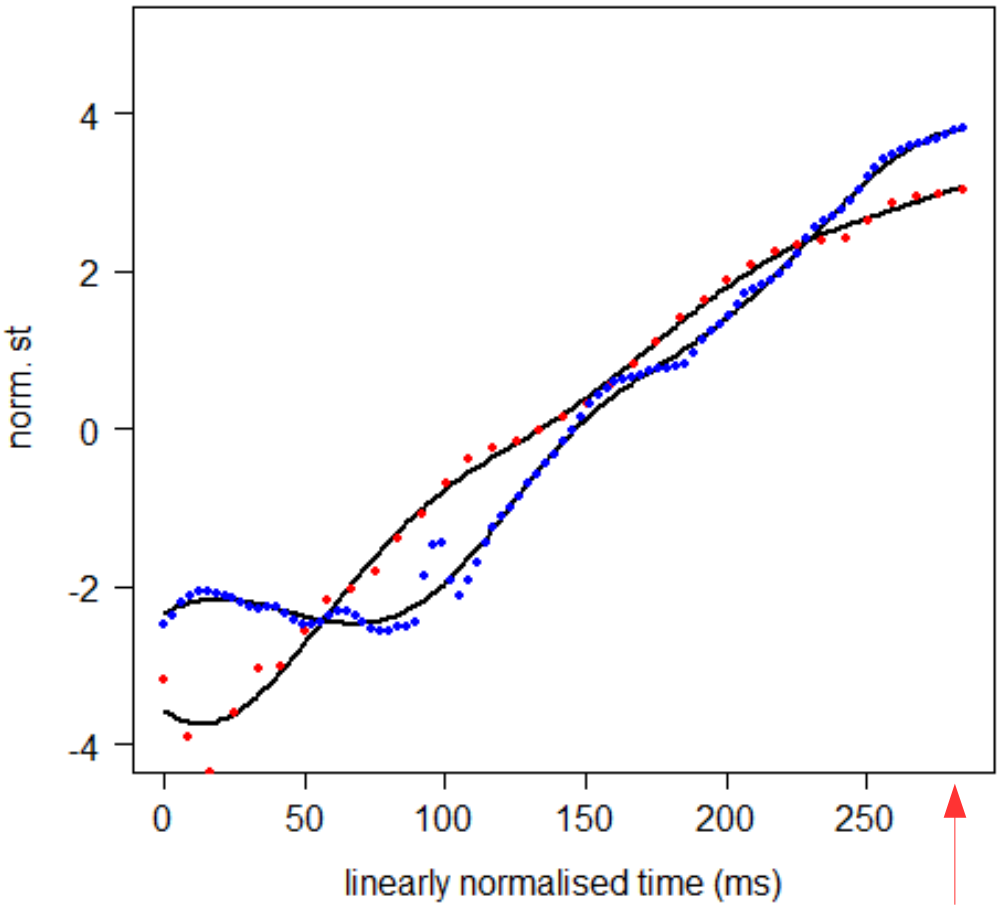

Linear time normalisation

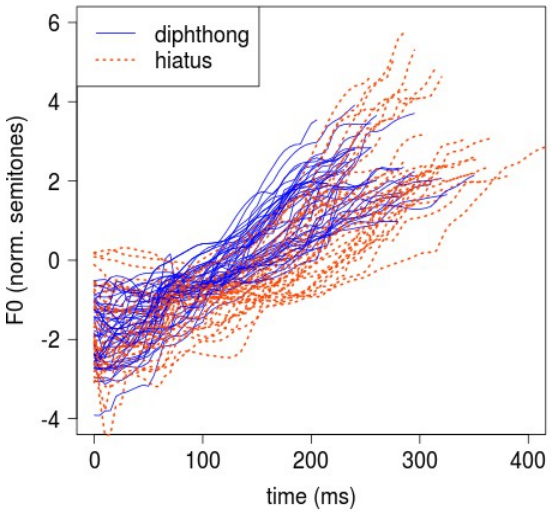

Interpolate to
the same
time interval

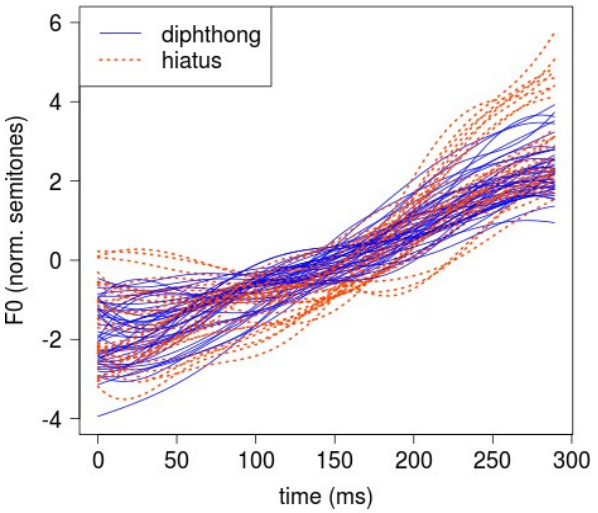

- We must use the same time interval
- This implies linear time normalisation
- Durations have to be reintroduced at the end of the analysis

Road map

CURVES

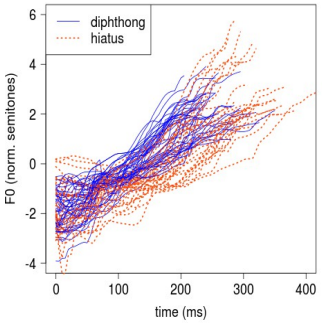

Interpolate to
<u>the same</u>
time interval

- Data driven

Dimensionality
reduction tool

- Few parameters
- Interpretable

NUMBERS

LMER

Introducing Functional PCA

Vectors

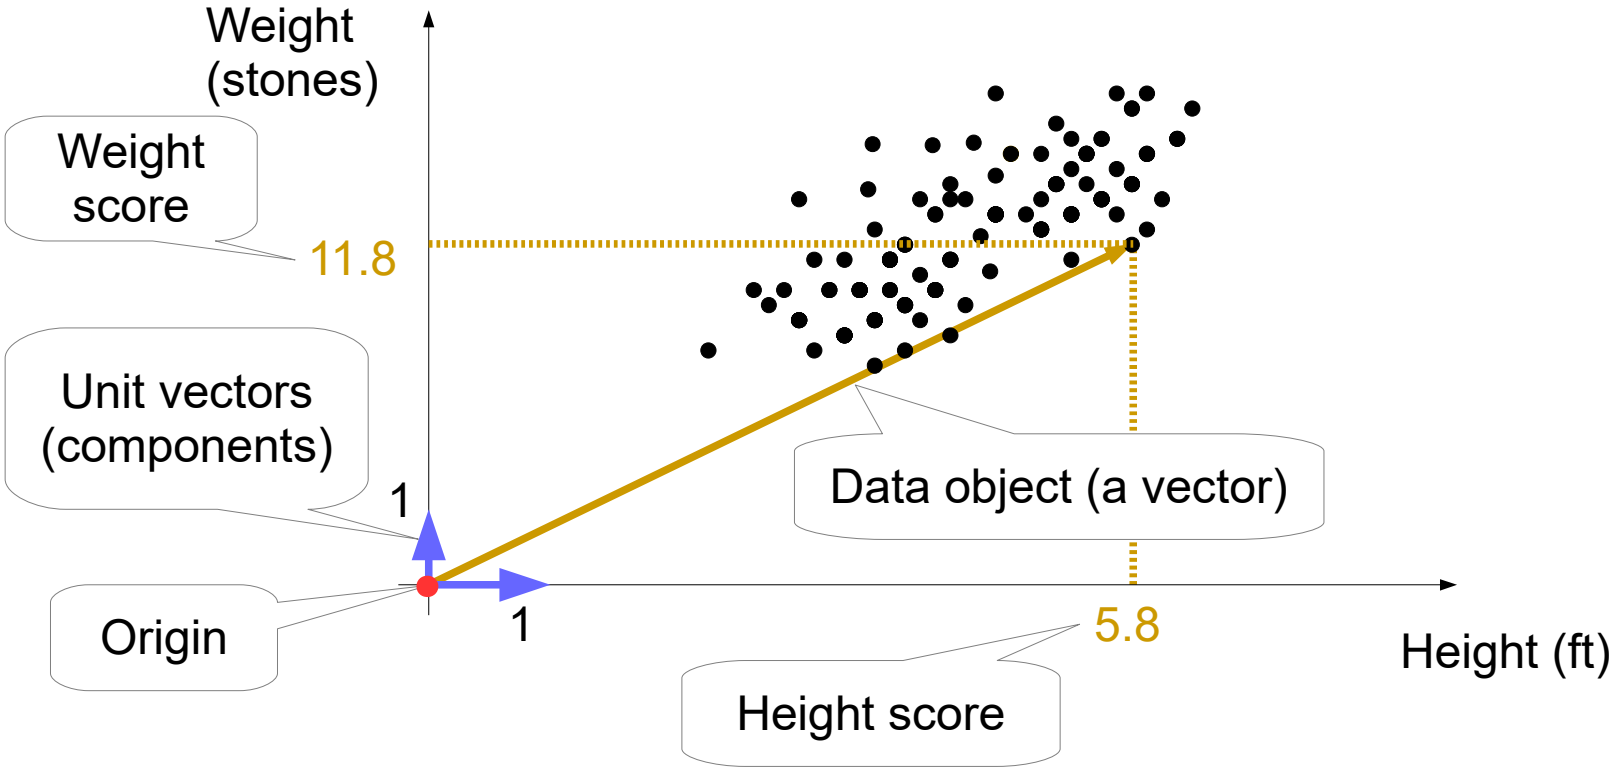

- Data objects and components are vectors
- From scores (numbers) we can reconstruct data objects (vectors)

Ordinary PCA

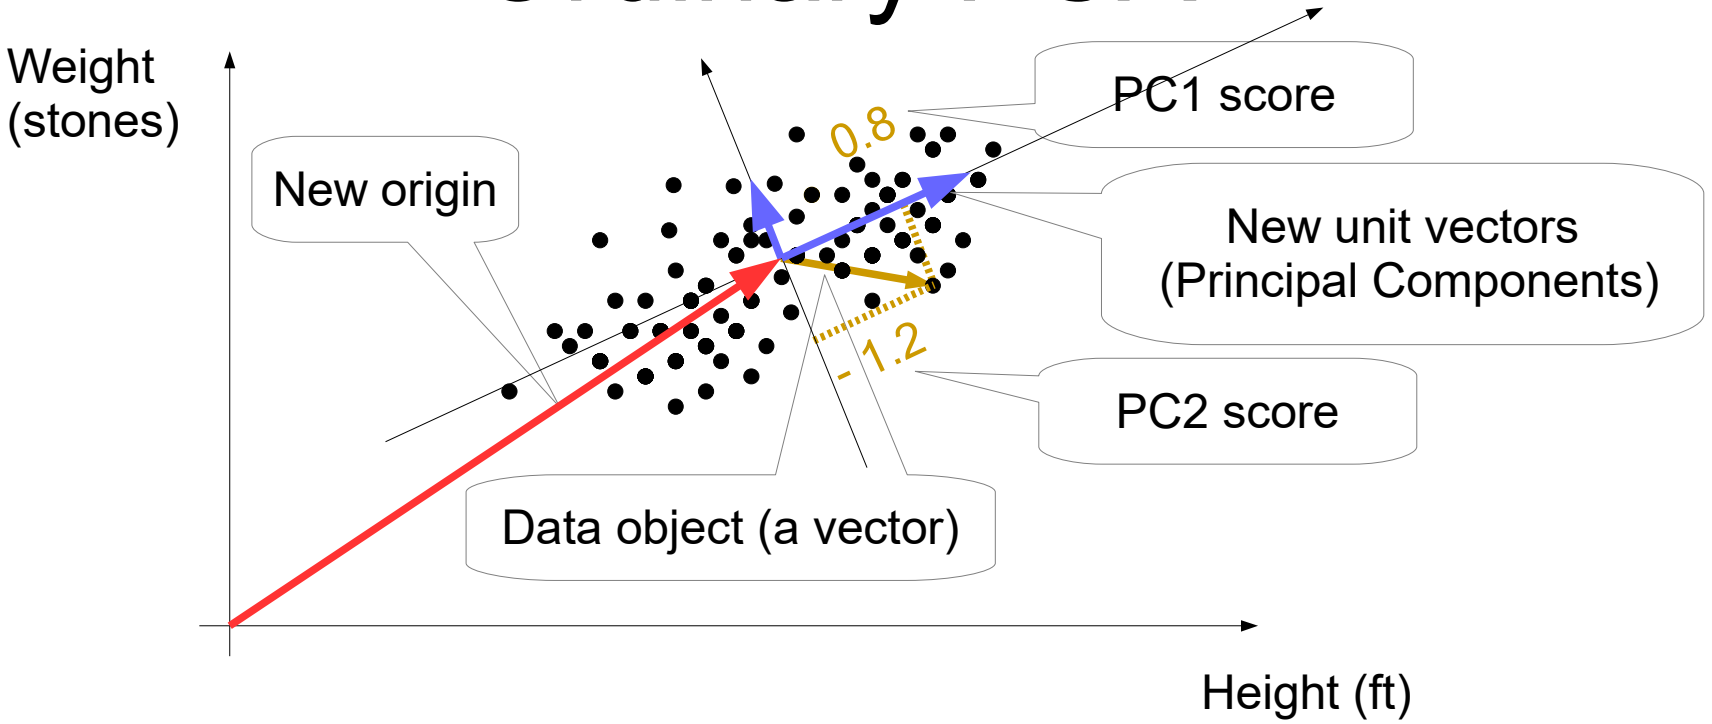

- PCA computes new origin and unit vectors which best suit the data
- From PC scores we can reconstruct data objects

Dimensionality reduction

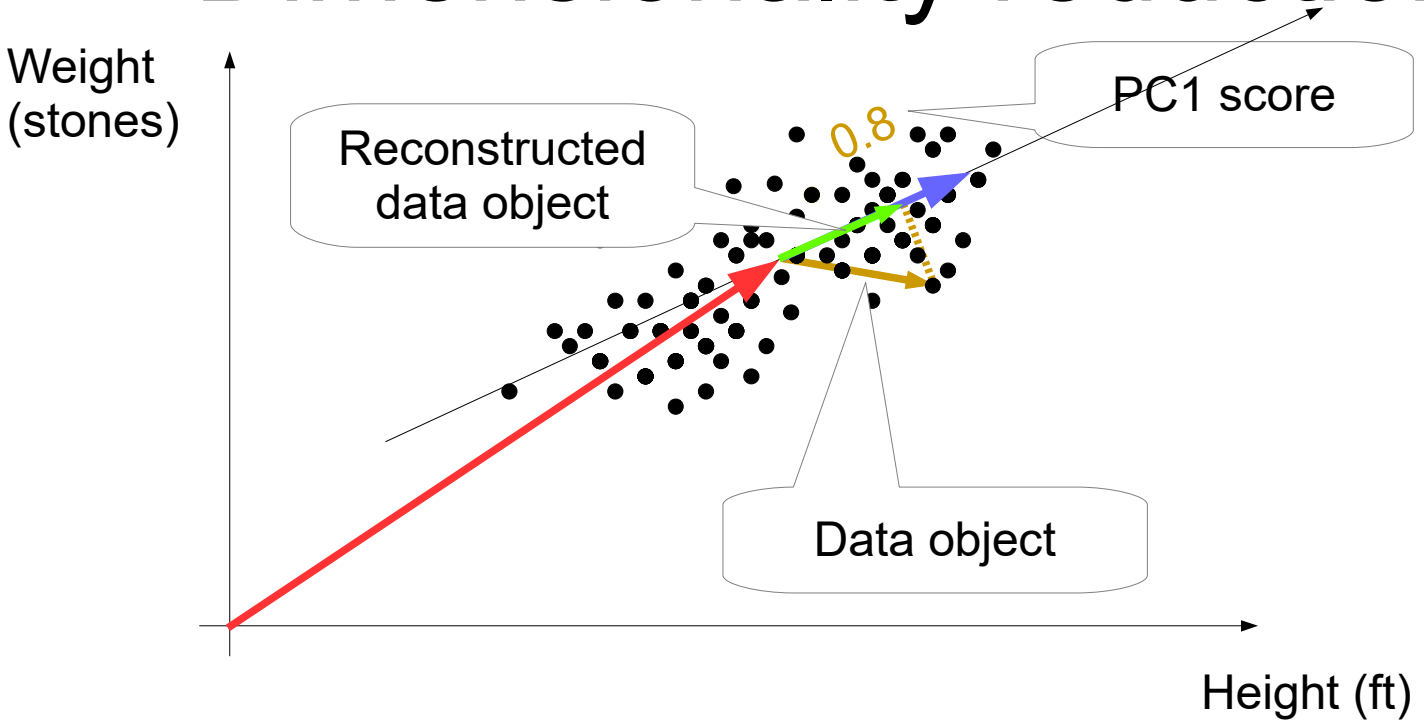

- We can use only part of the PCs
- This reduces the data dimensionality
- But introduces reconstruction errors too

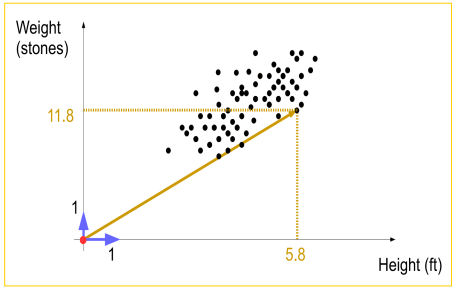

Functions (curves)

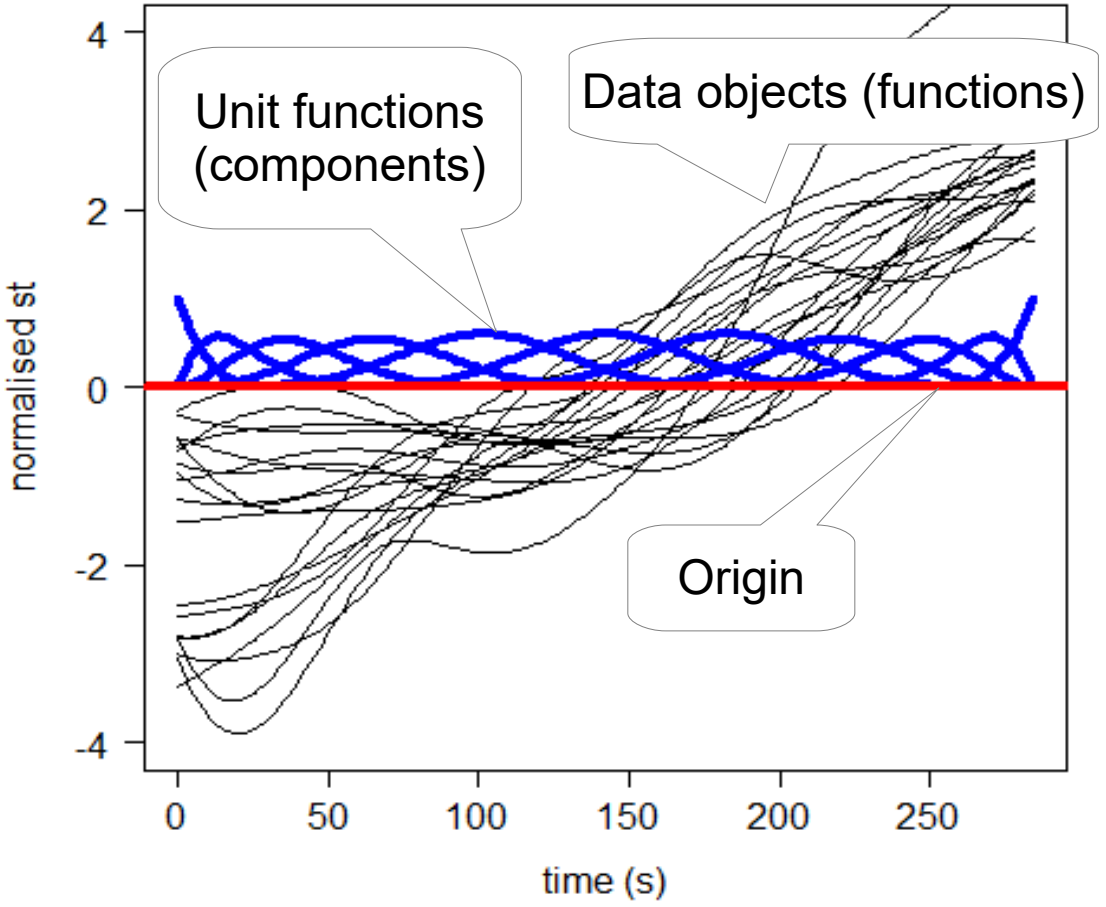

- Origin, components and data objects are functions
- Origin is a flat line
- Components are 11 B-spline curves

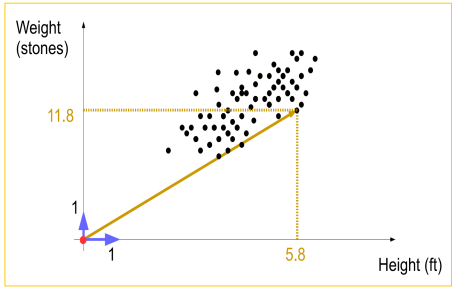

Functions (curves)

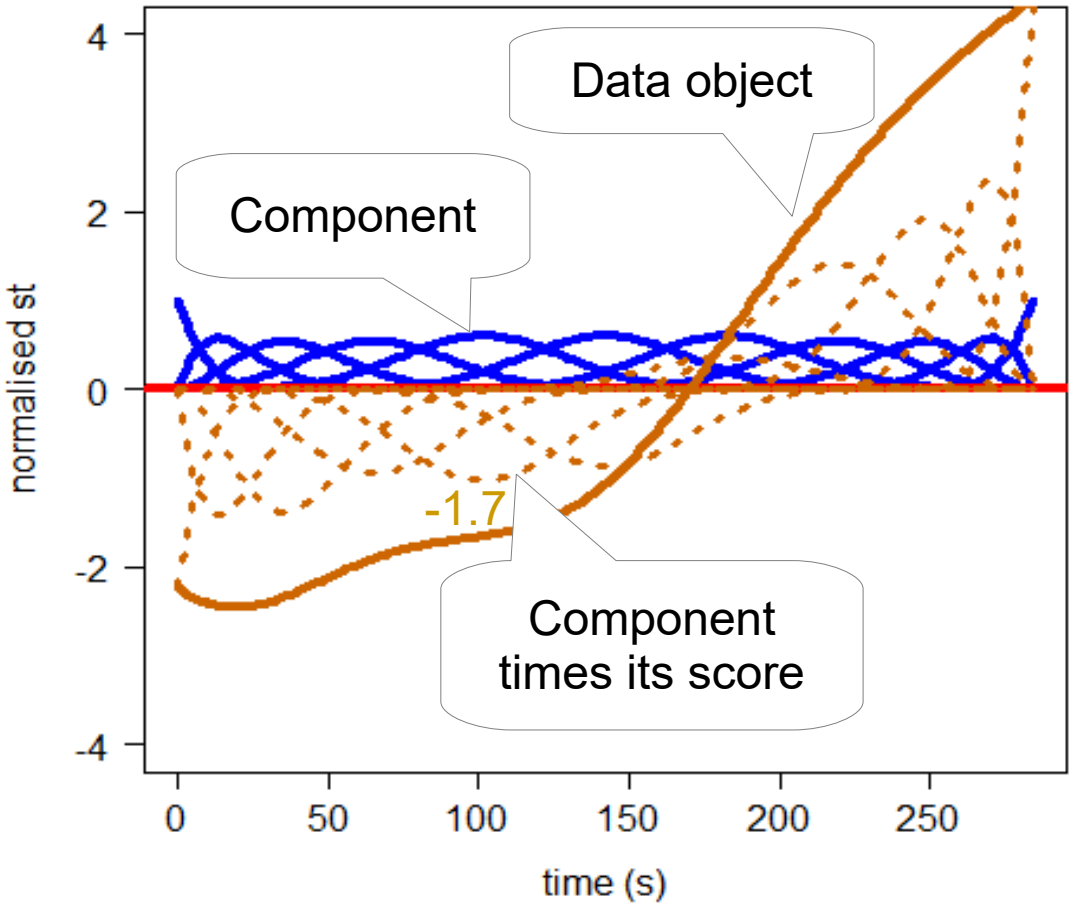

- Each of the 11 components is multiplied by a score
- These are summed together to obtain a data object

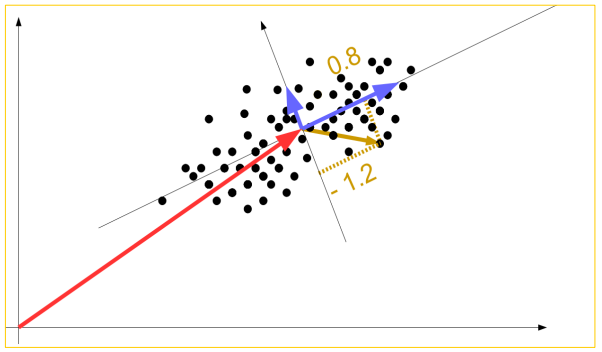

Functional PCA

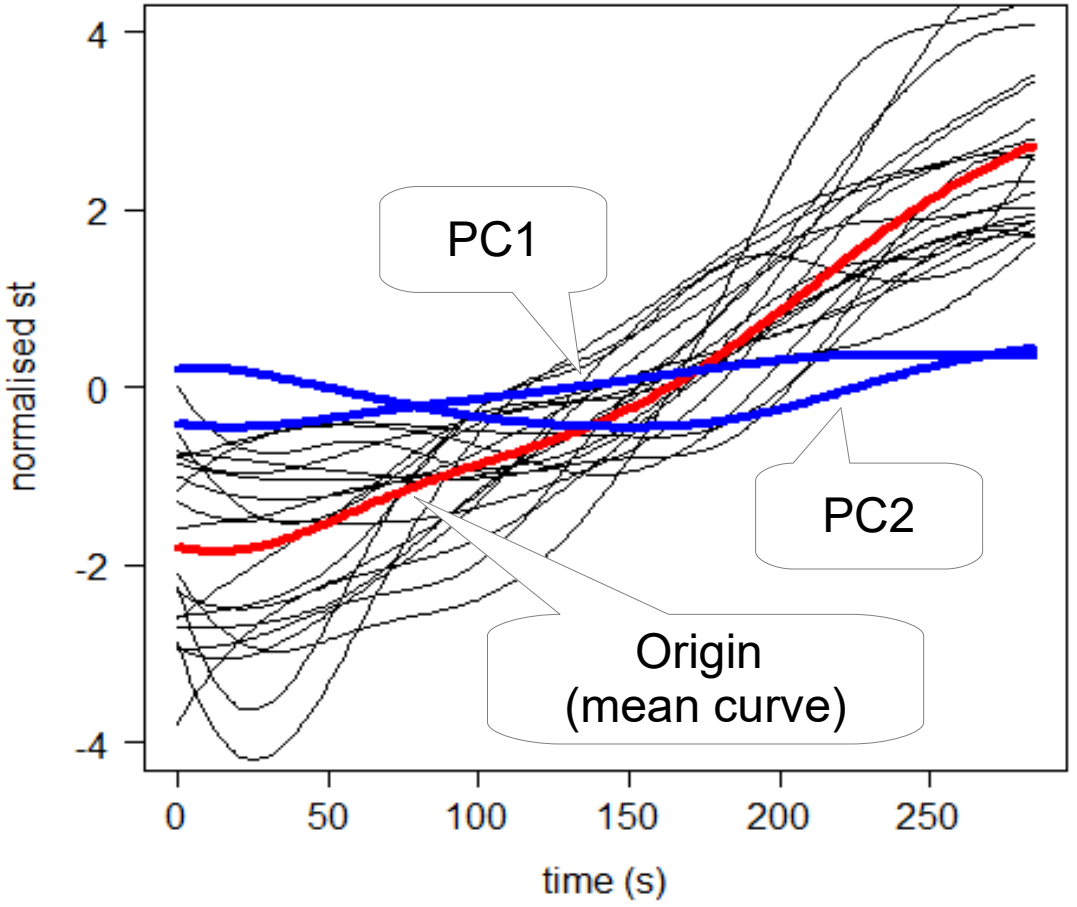

- FPCA computes new origin and component functions which best suit the data

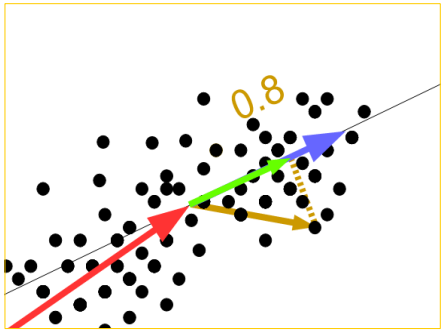

Functional PCA

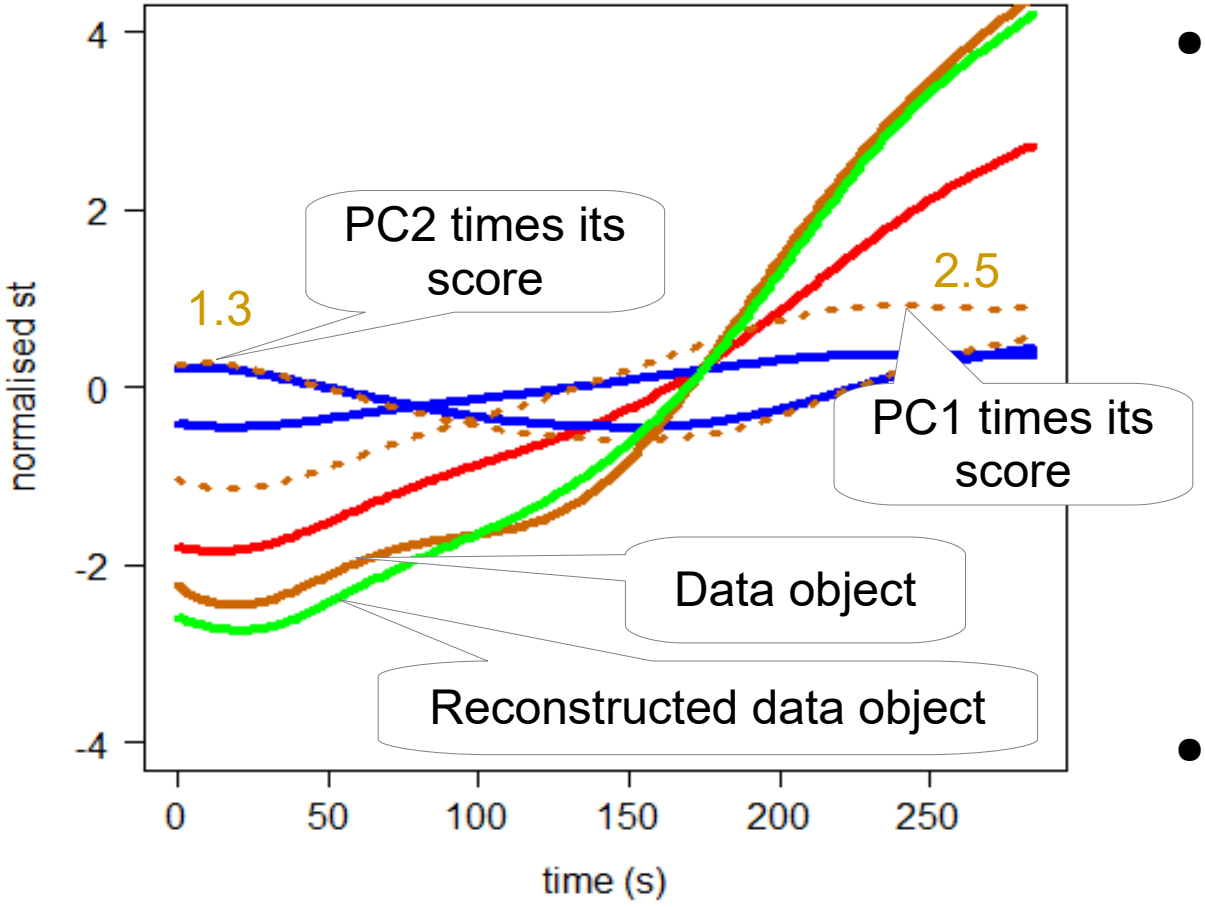

- The sum of origin (mean) curve + PCs times their scores gives an approx reconstruction of the original curve
- Dimensions from 11 (B-splines) down to 2 (PCs)

Functional PCs

$$f(t) \approx \mu(t) + s_1 \cdot PC1(t) + s_2 \cdot PC2(t) + \cdots$$

PC1 score

PC2 score

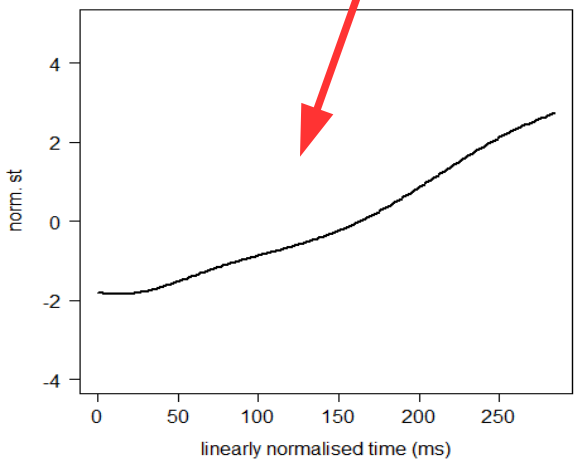

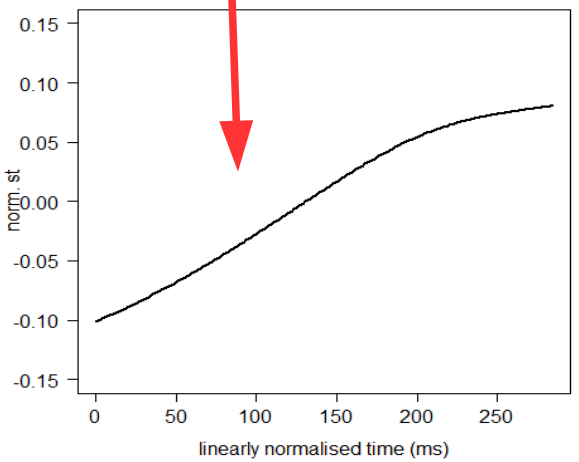

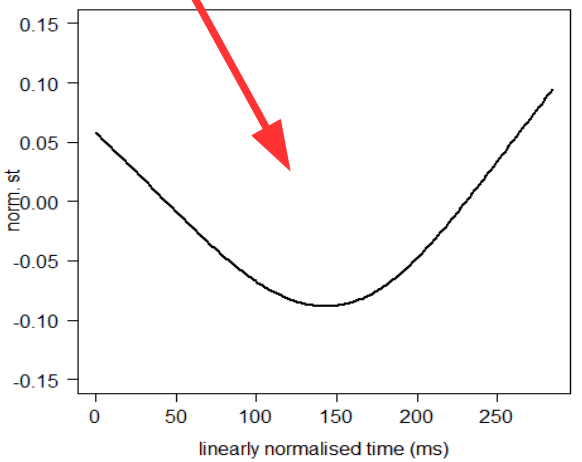

Curve reconstruction

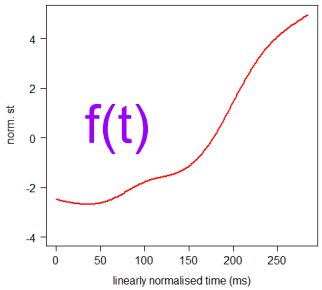

≈

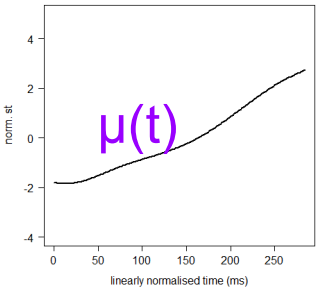

+ 16.5 *

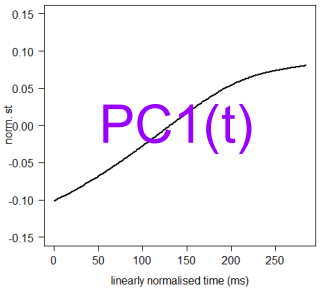

+ 10.8 *

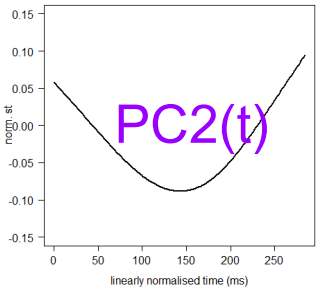

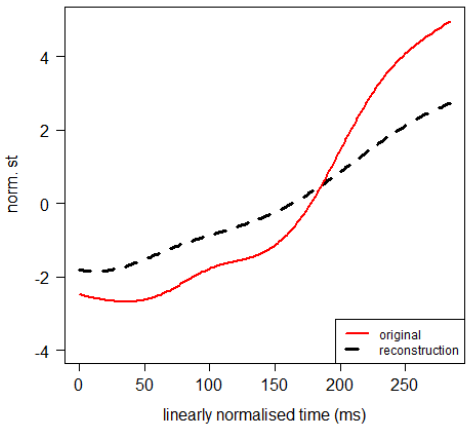

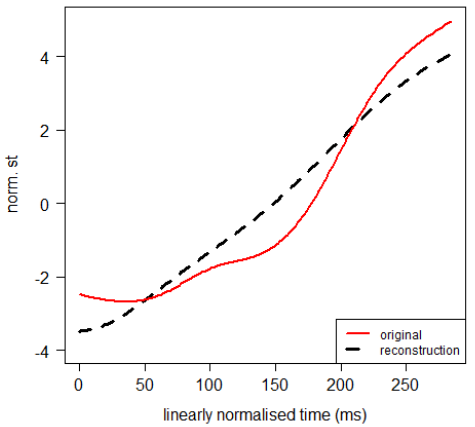

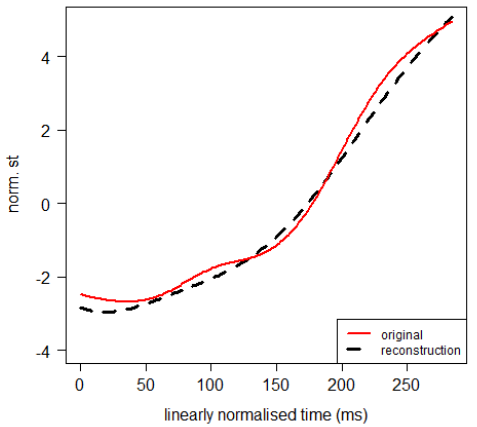

PC1 scores

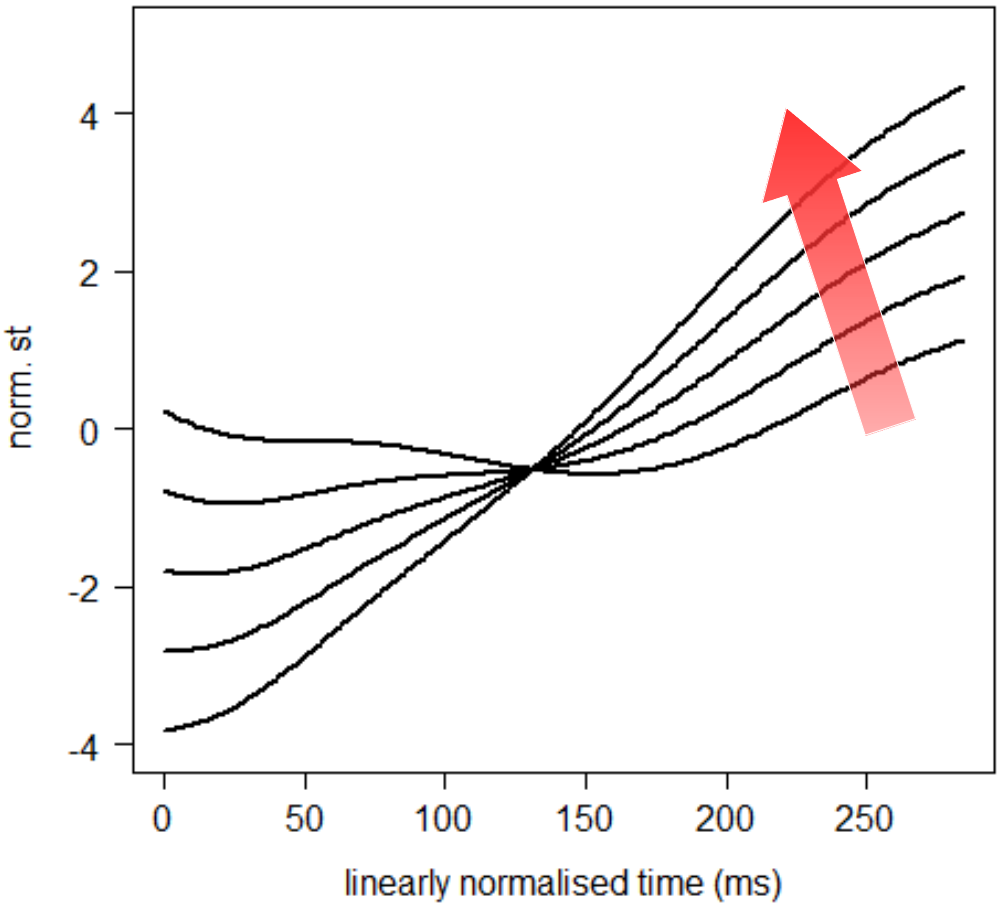

μ(t) + **20** * PC1(t)
μ(t) + **10** * PC1(t)
μ(t) + **0** * PC1(t)
μ(t) - **10** * PC1(t)
μ(t) - **20** * PC1(t)

PC2 scores

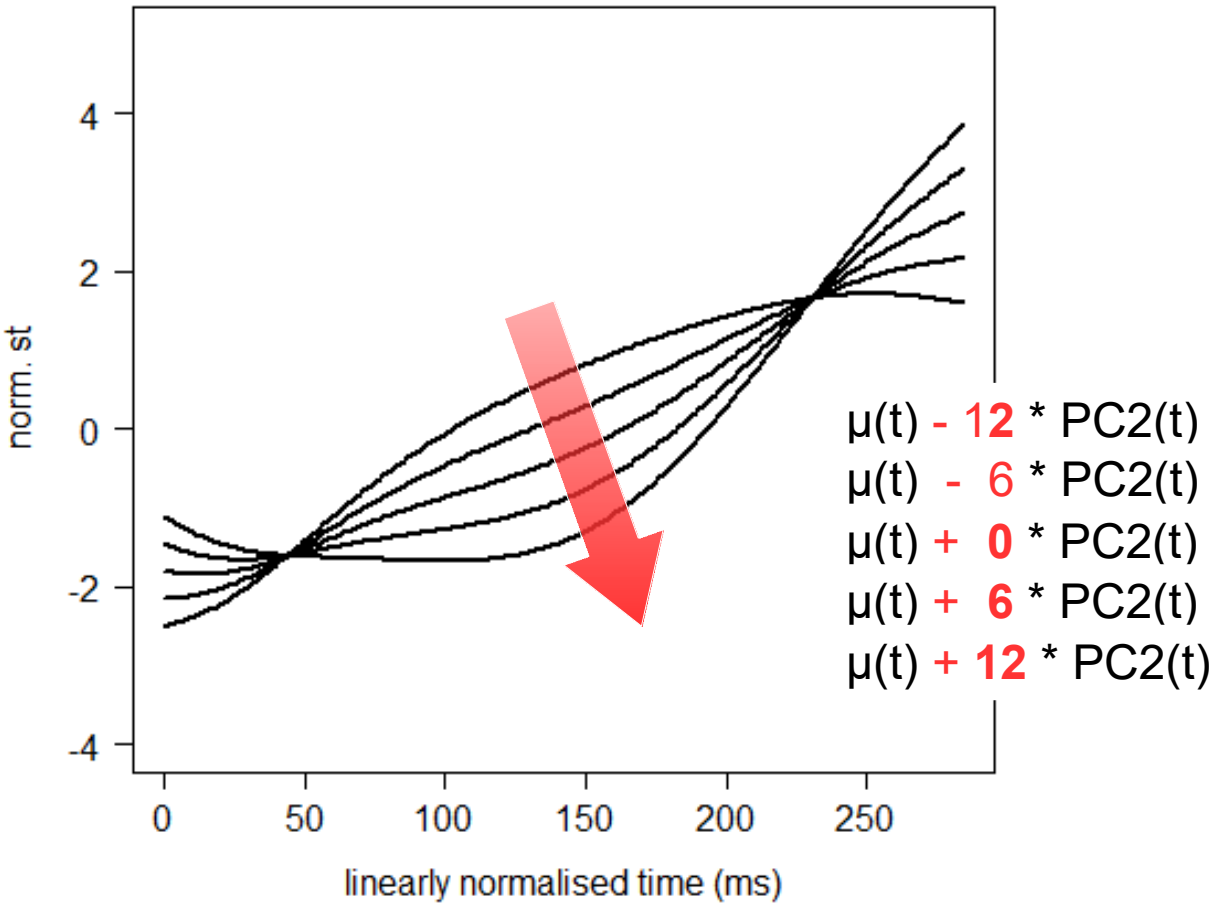
4
2
0
-2
-4
norm. st
0
50
100
150
200
250
linearly normalised time (ms)
μ(t) - 12 * PC2(t)
μ(t) - 6 * PC2(t)
μ(t) + 0 * PC2(t)
μ(t) + 6 * PC2(t)
μ(t) + 12 * PC2(t)

Curve parametrisation

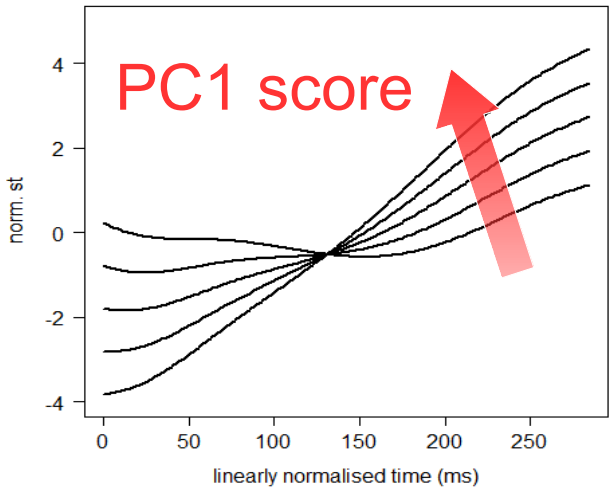

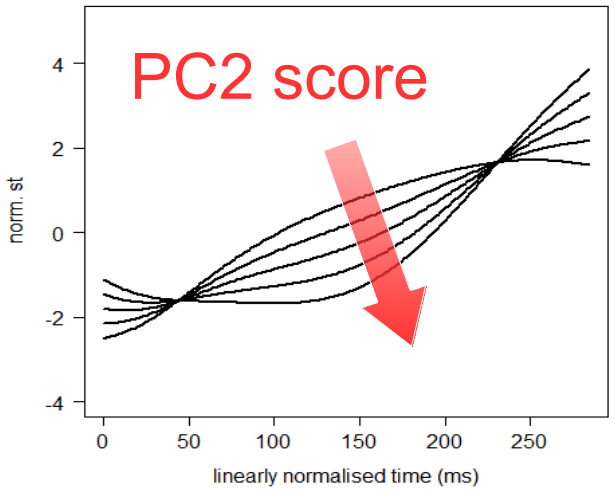

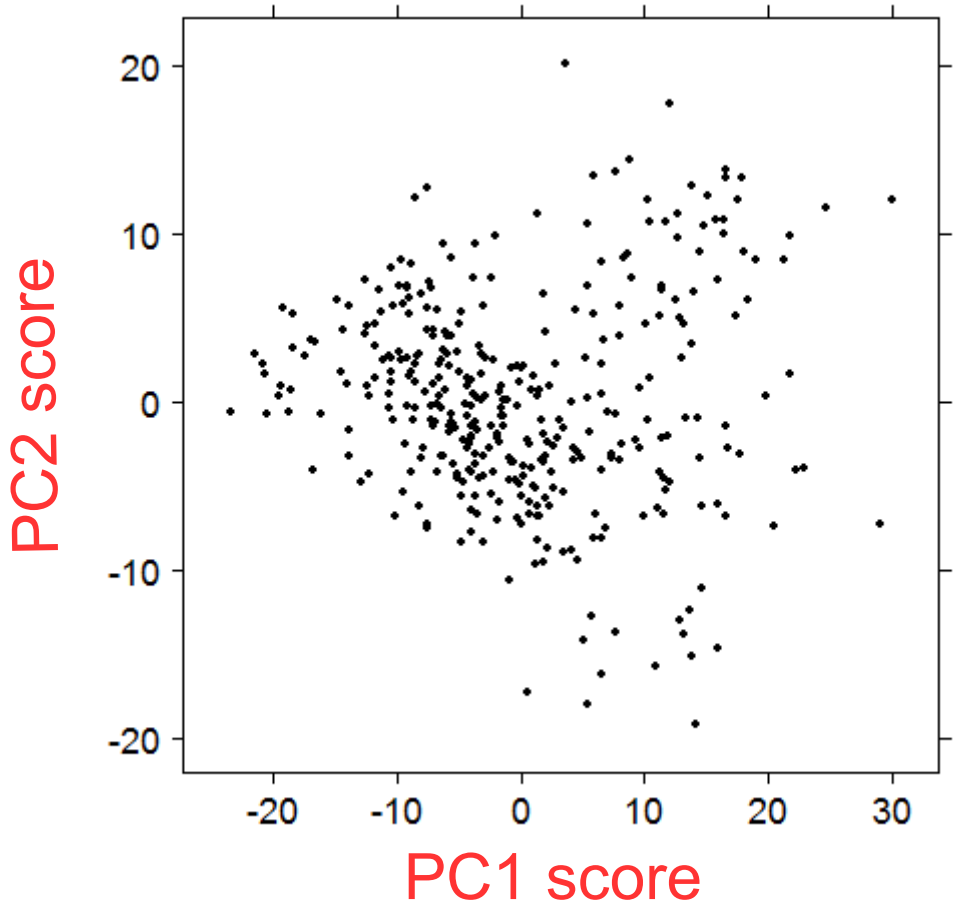

Curve parametrisation

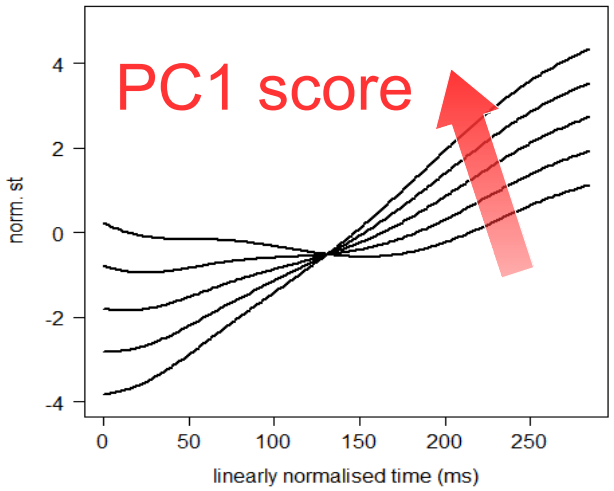

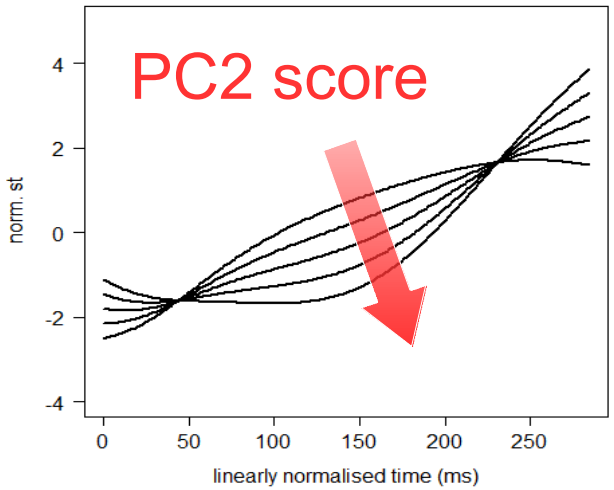

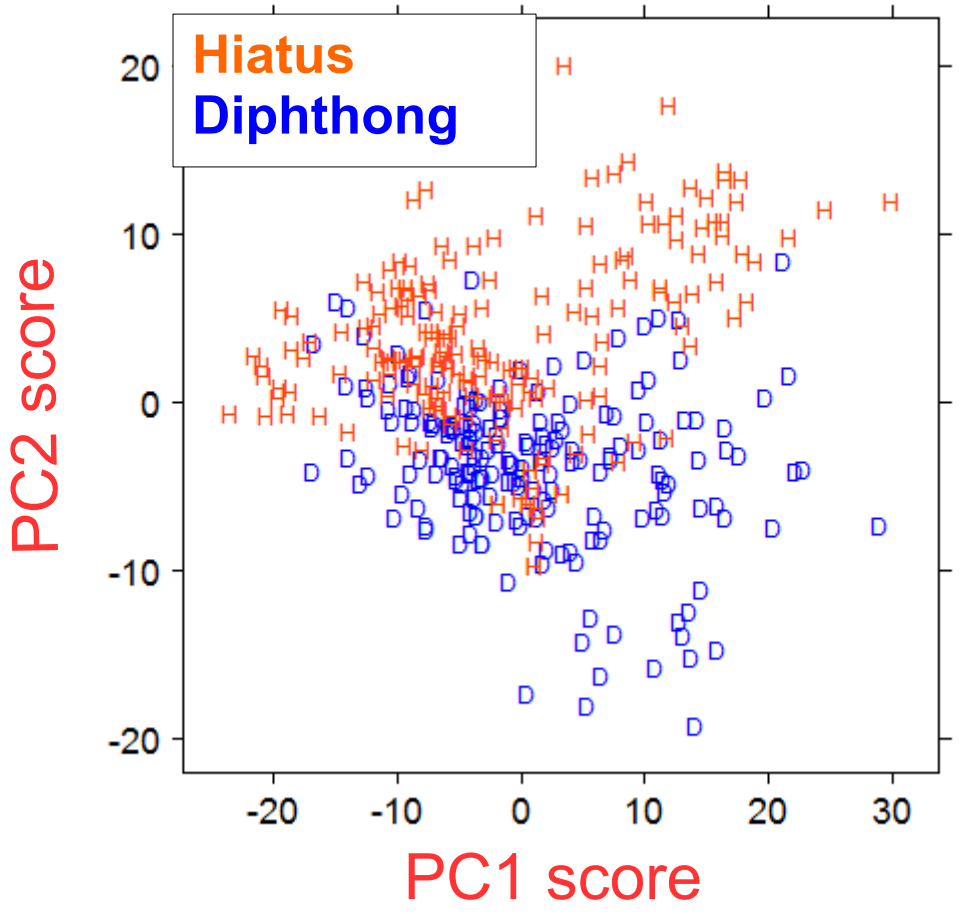

Multidimensional signals

Formants

2D CURVES

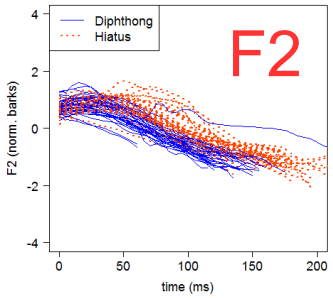

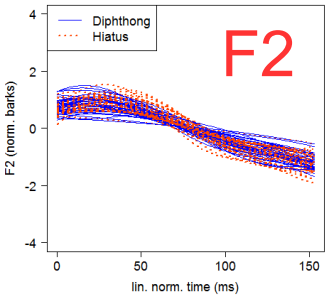

FPCA

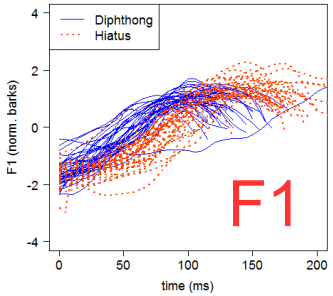

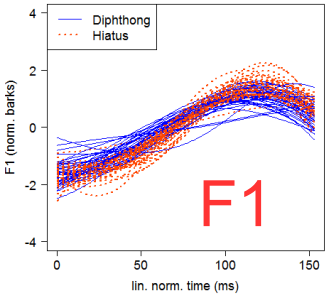

FPCA

NUMBERS

LMER

Formants

2D CURVES

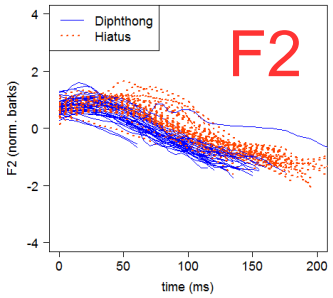

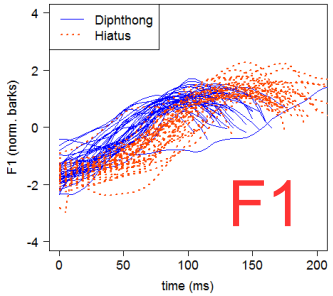

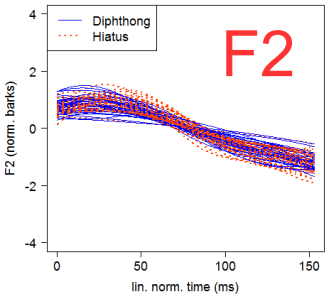

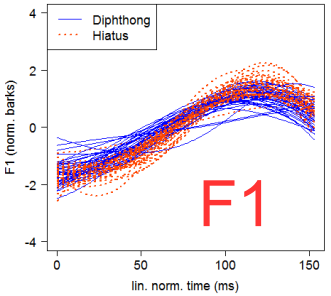

2D
FPCA

NUMBERS

LMER

Trick

2D CURVES

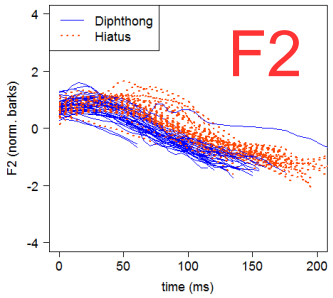

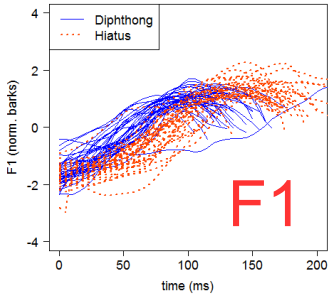

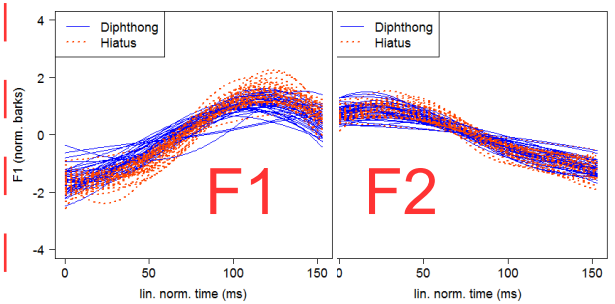

FPCA

NUMBERS

LMER

PC1 scores

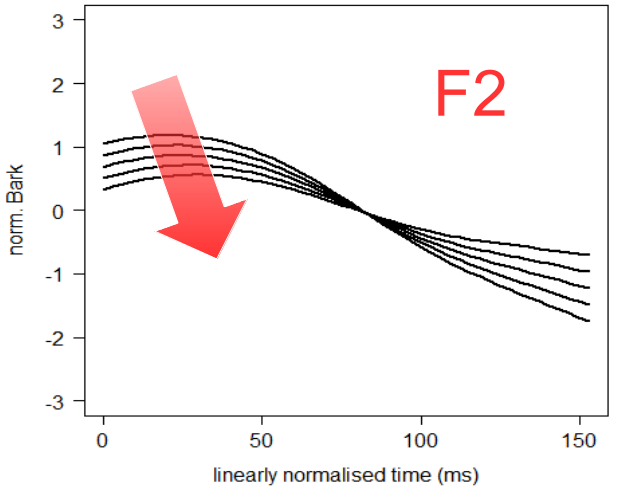

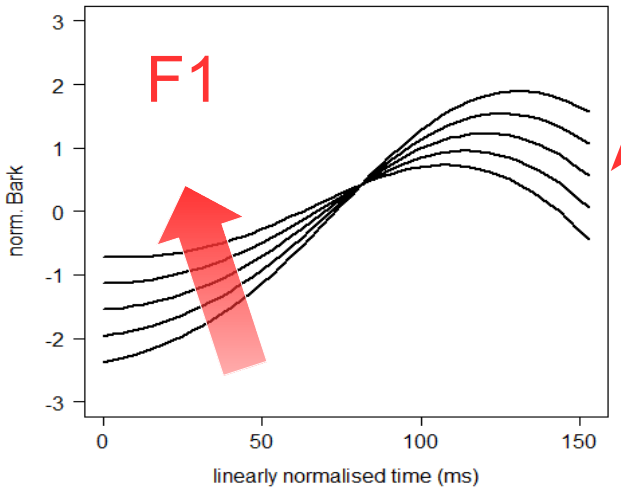

μ(t) + **8** * PC1(t)
μ(t) + **4** * PC1(t)
μ(t) + **0** * PC1(t)
μ(t) - **4** * PC1(t)
μ(t) - **8** * PC1(t)

PC2 scores

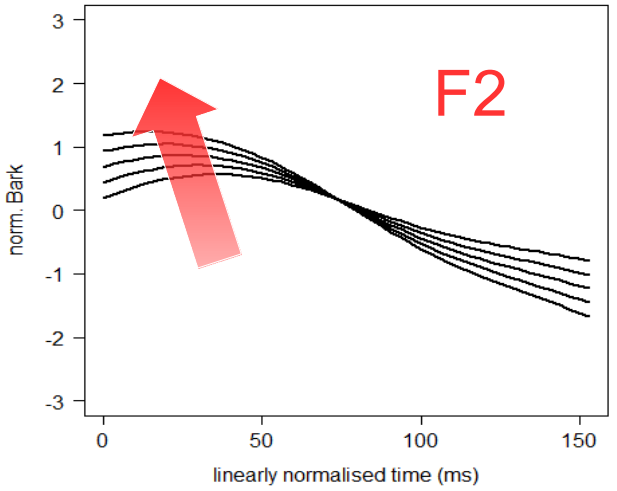

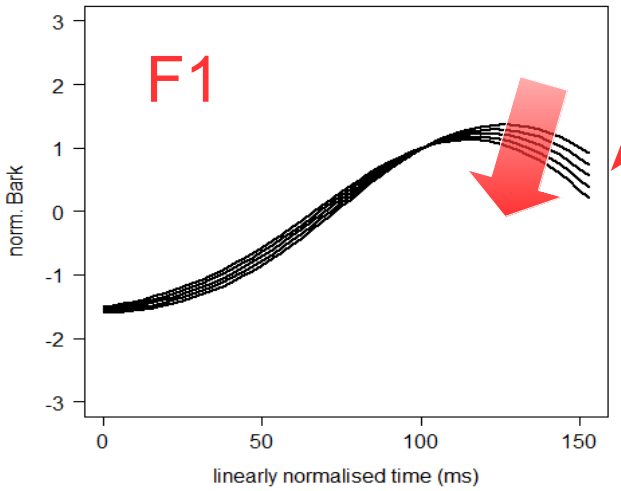

μ(t) + **4** * PC1(t)
μ(t) + **2** * PC1(t)
μ(t) + **0** * PC1(t)
μ(t) - **2** * PC1(t)
μ(t) - **4** * PC1(t)

2D curve parametrisation

PC1 score

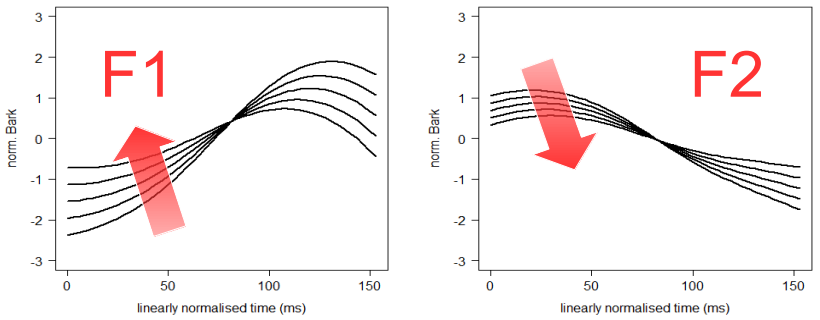

PC2 score

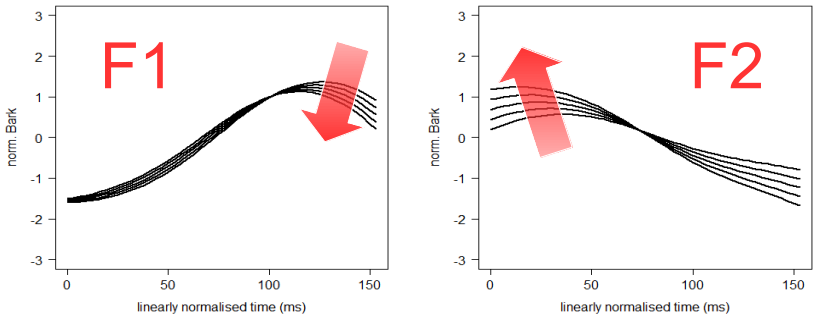

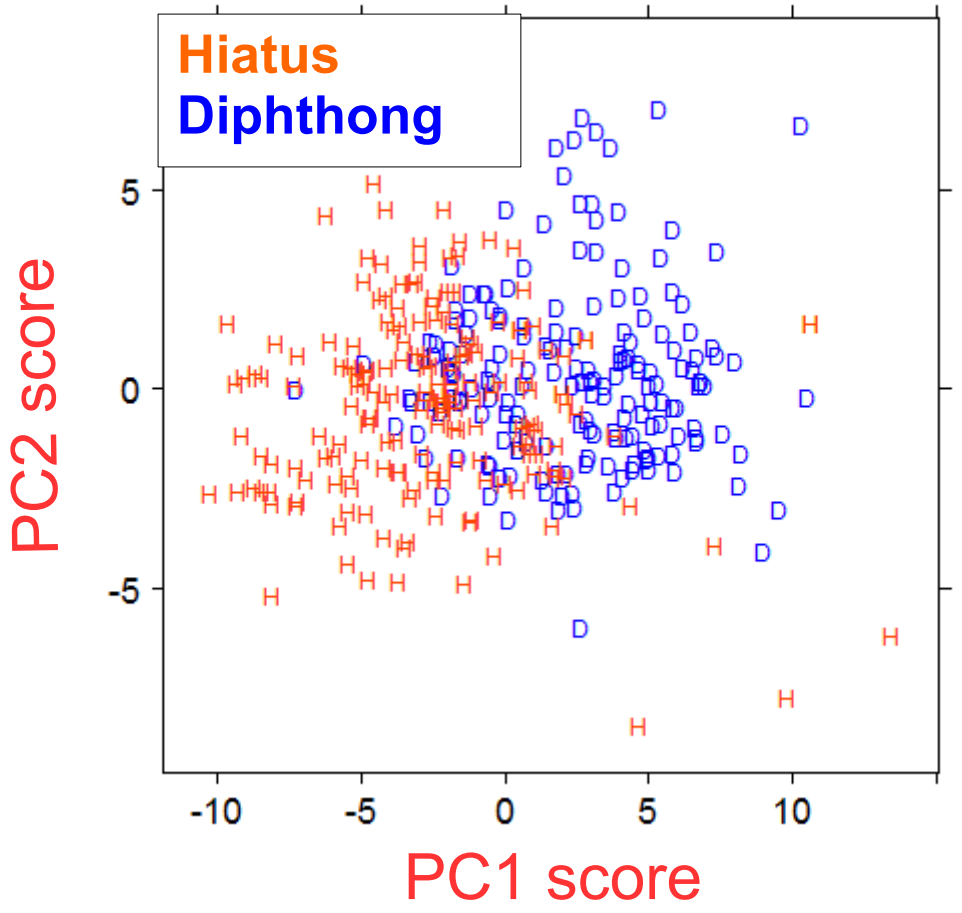

Long signals

Many segments

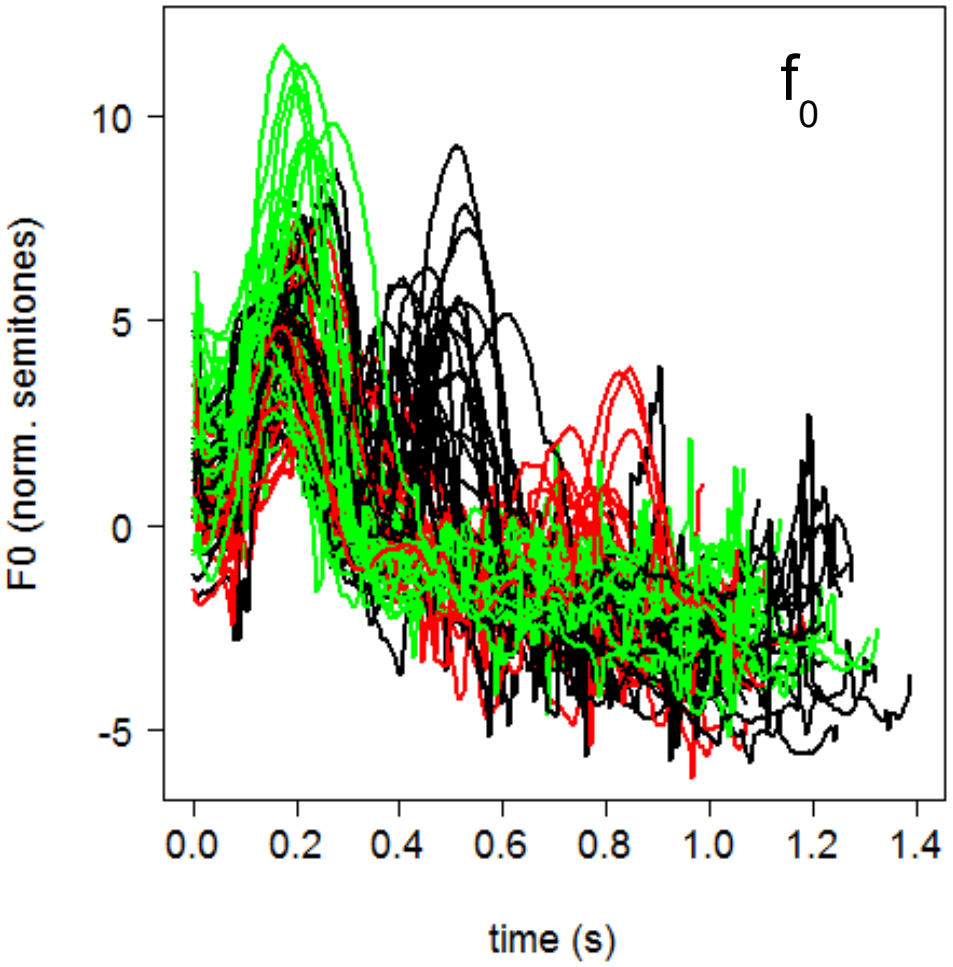

- Narrow focus in Neapolitan Italian
- Focus on

  Subject, **Verb** or **Prop. Phrase**

  Danilo vola da Roma

  (*Danilo flies from Rome*)
- 8 CV syllables

  first C was excluded (too short)

  VCVCV CVCV CV CVCV

… 15 segments!

Linear time normalisation

BEFORE

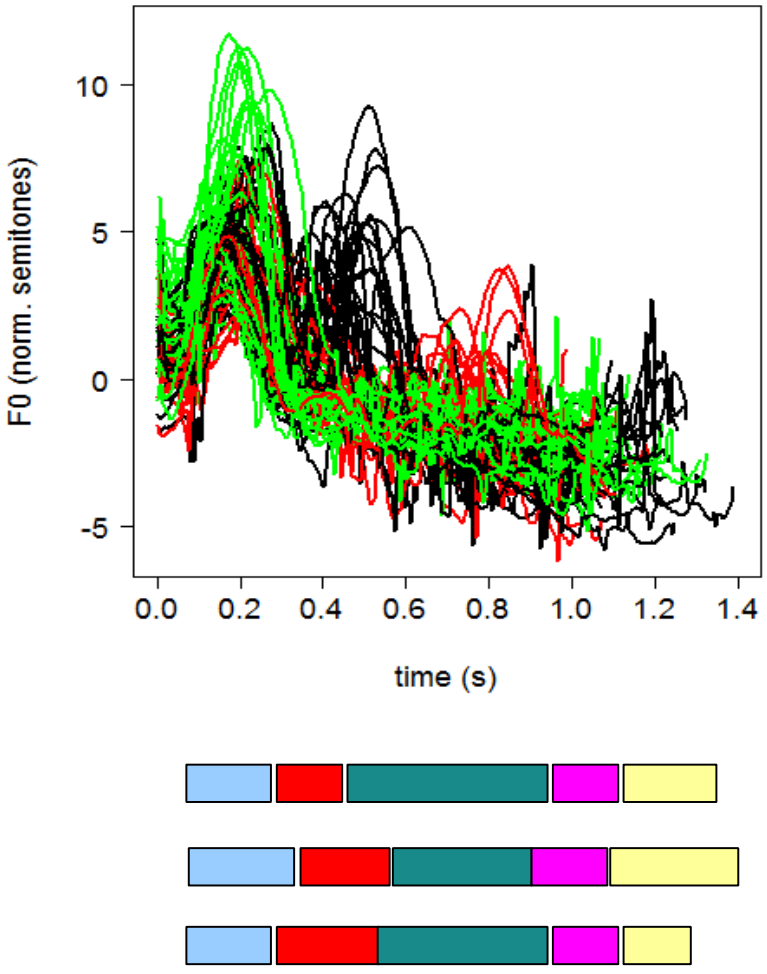

AFTER

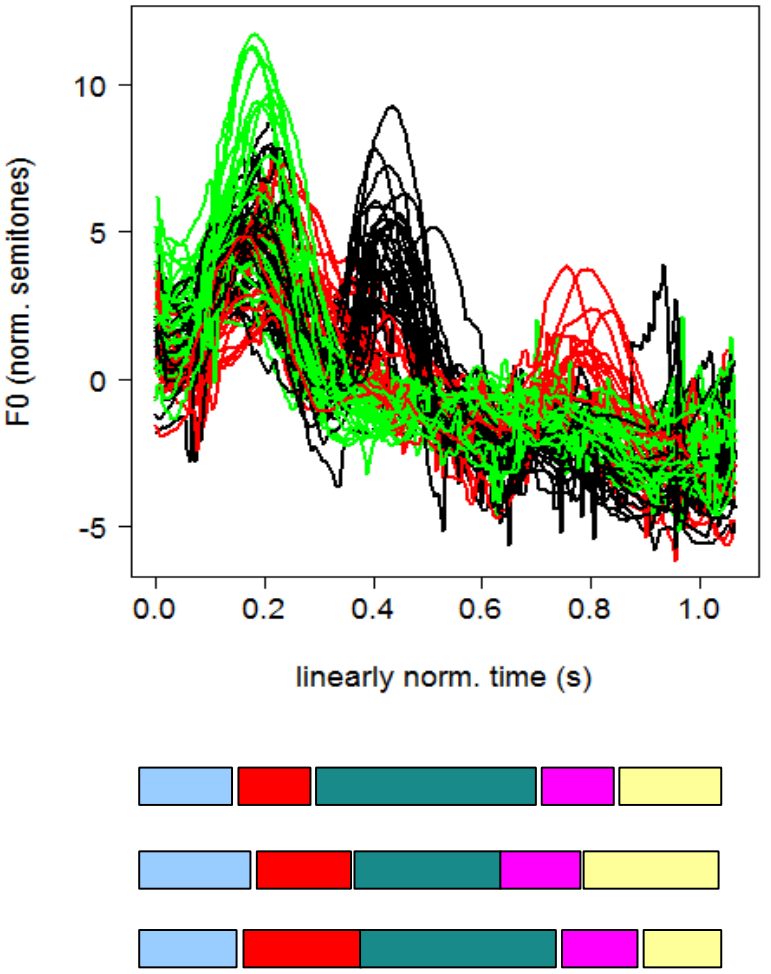

Landmark registration

BEFORE

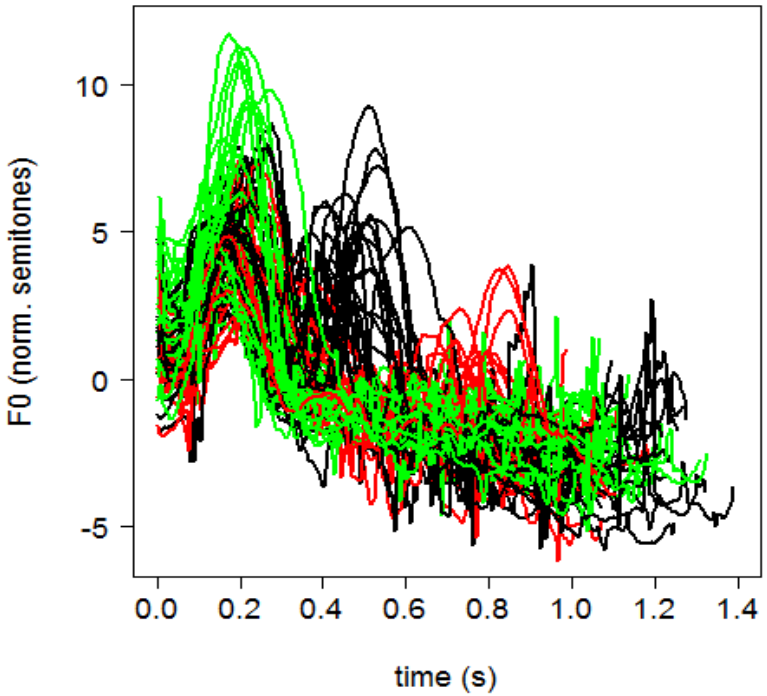

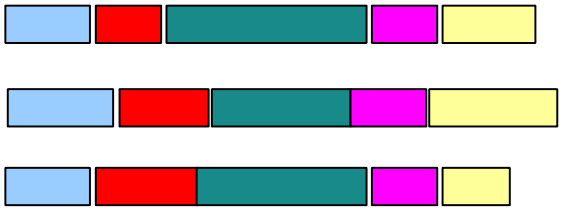

AFTER

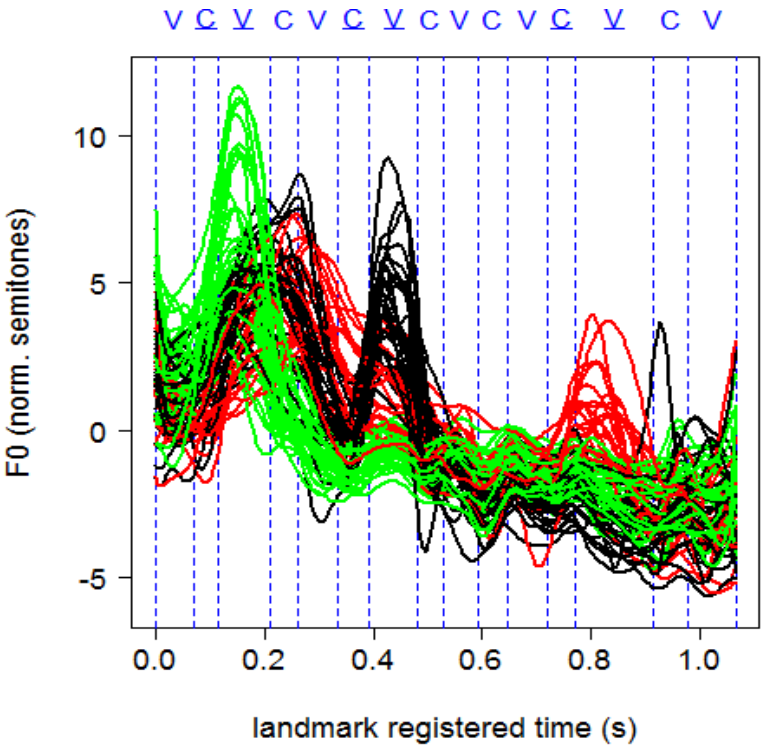

Using landmark registration

CURVES

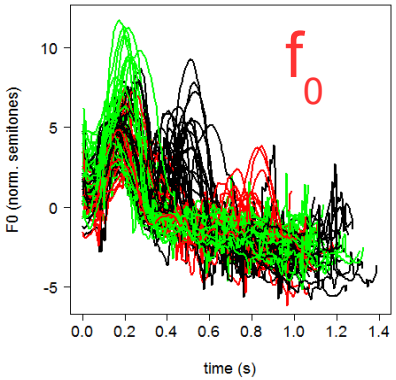

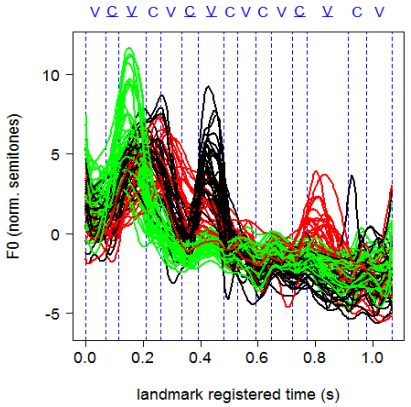

segment durations

d1	d2	...	d15
...	...	...	...
...	...	...	...
...	...	...	...

FPCA

PCA

NUMBERS

LMER

Using landmark registration

CURVES

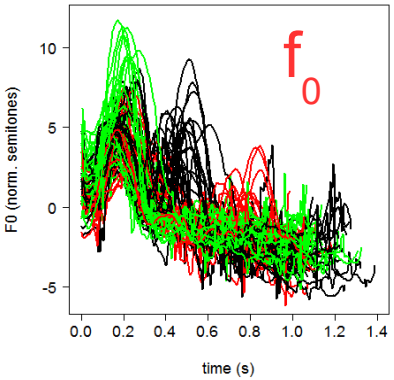

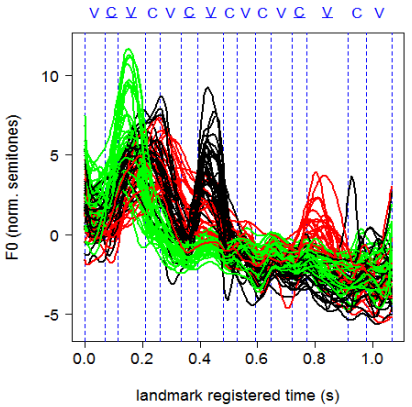

segment durations

d1	d2	...	d15
...	...	...	...
...	...	...	...
...	...	...	...

FPCA

NUMBERS

LMER

Inside landmark registration

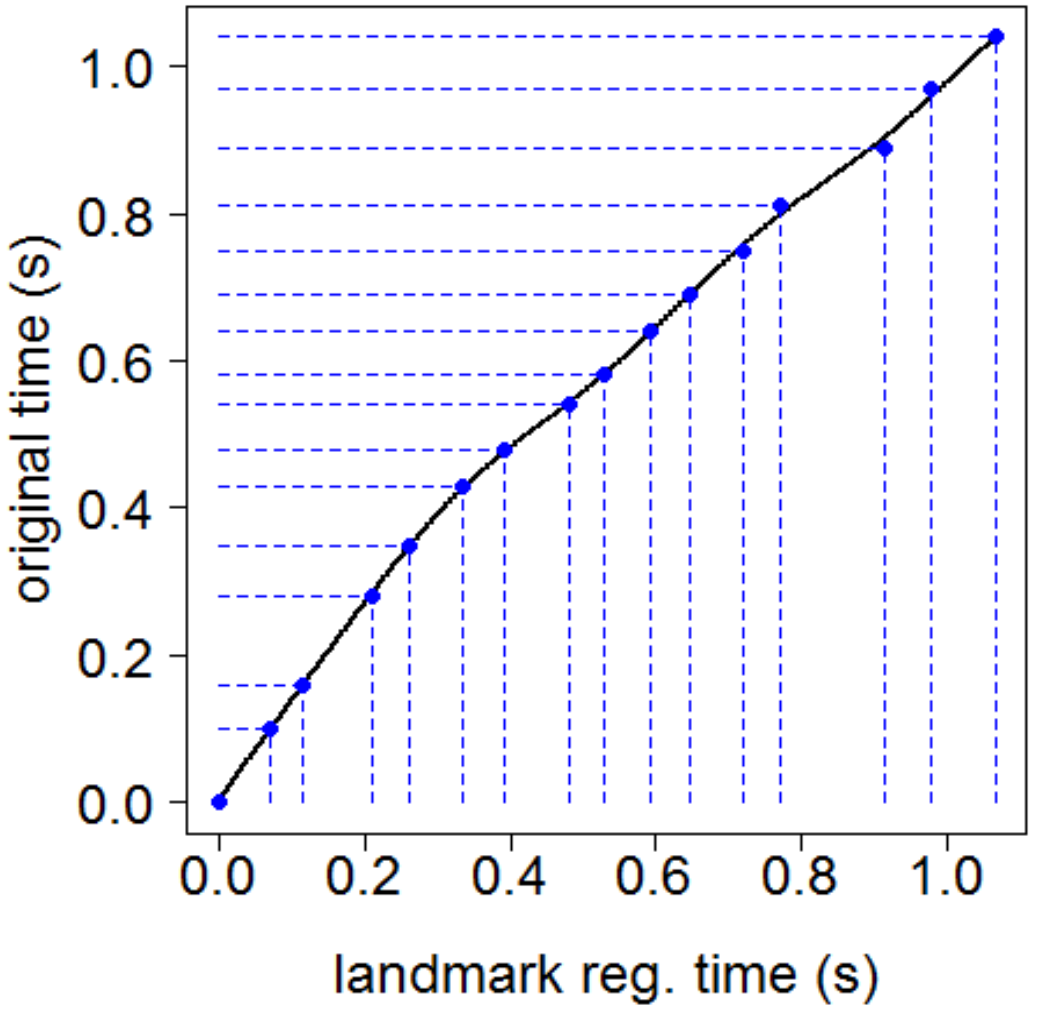

h(t)

Relative log rate

1

h(t)

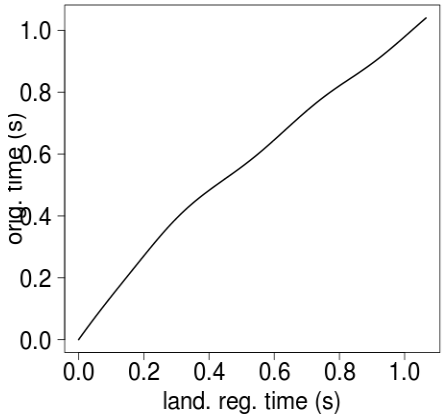

2

dh(t)/dt

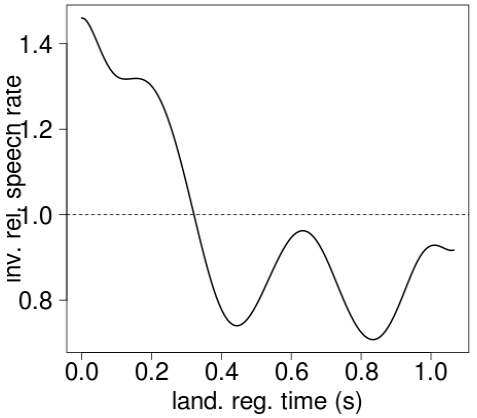

REVERSIBLE!

1 ⟷ 4

3

- dh(t)/dt

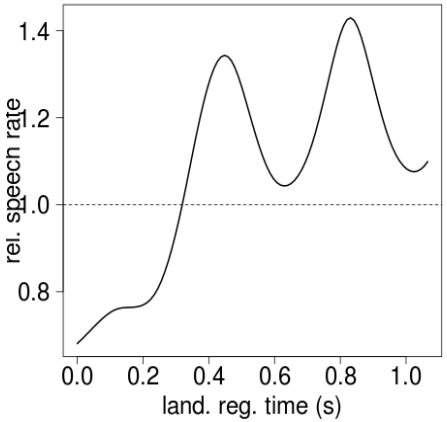

4

- log dh(t)/dt

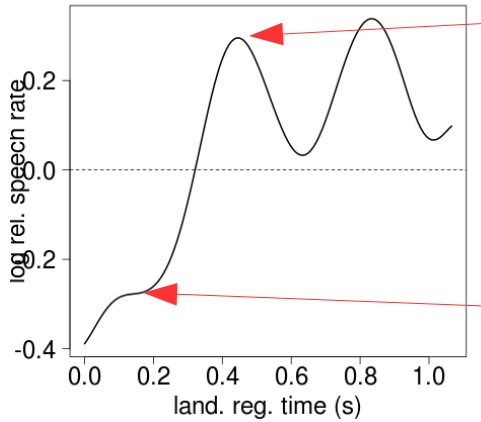

+ 0.25 → duration / 1.28

0 → same duration

- 0.25 → duration * 1.28

Using log rates

CURVES

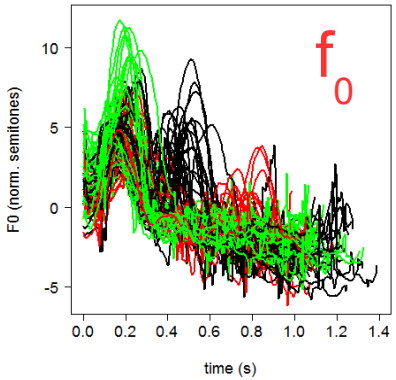

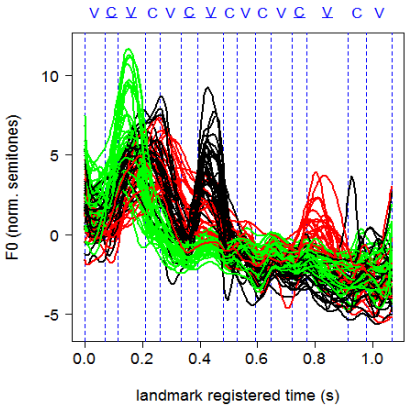

log rates

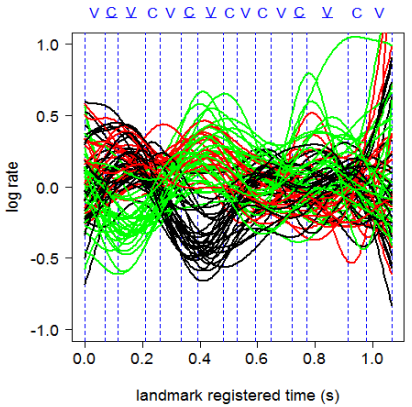

2D FPCA

NUMBERS

LMER

PC1 scores

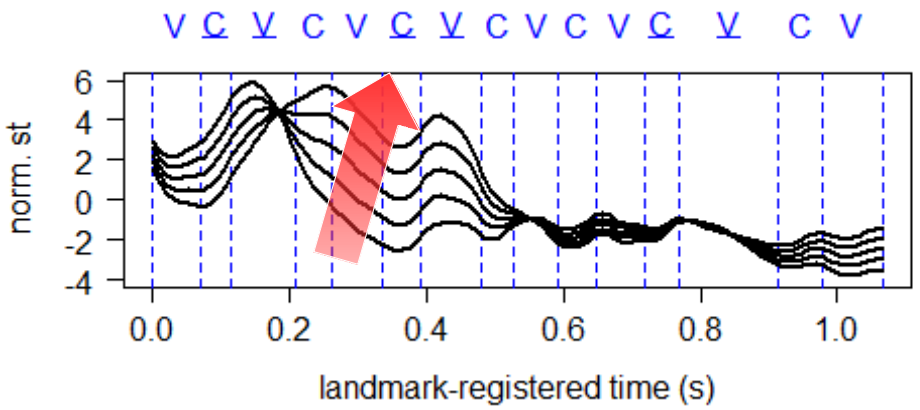

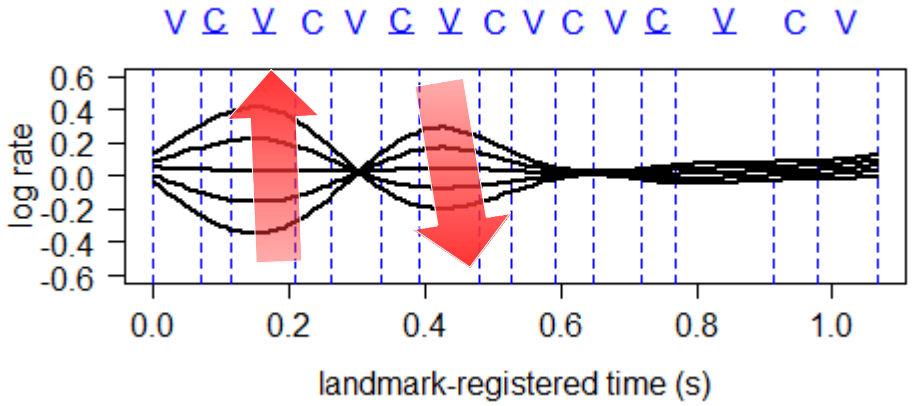

f_0

log rates

$\mu(t) + 2 * PC1(t)$
$\mu(t) + 1 * PC1(t)$
$\mu(t) + 0 * PC1(t)$
$\mu(t) - 1 * PC1(t)$
$\mu(t) - 2 * PC1(t)$

PC1 scores

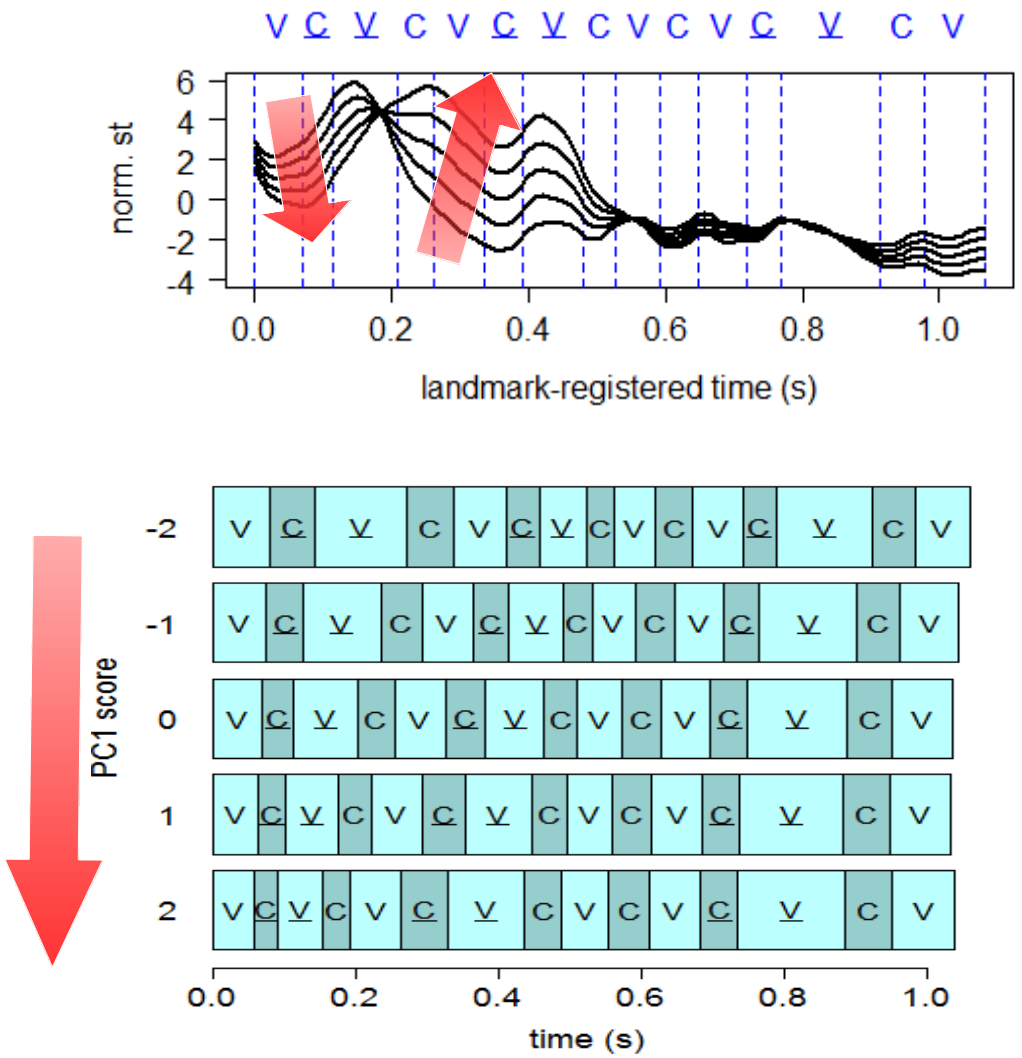

f_0

$\mu(t) + 2 * PC1(t)$
$\mu(t) + 1 * PC1(t)$
$\mu(t) + 0 * PC1(t)$
$\mu(t) - 1 * PC1(t)$
$\mu(t) - 2 * PC1(t)$

segment durations

PC2 scores

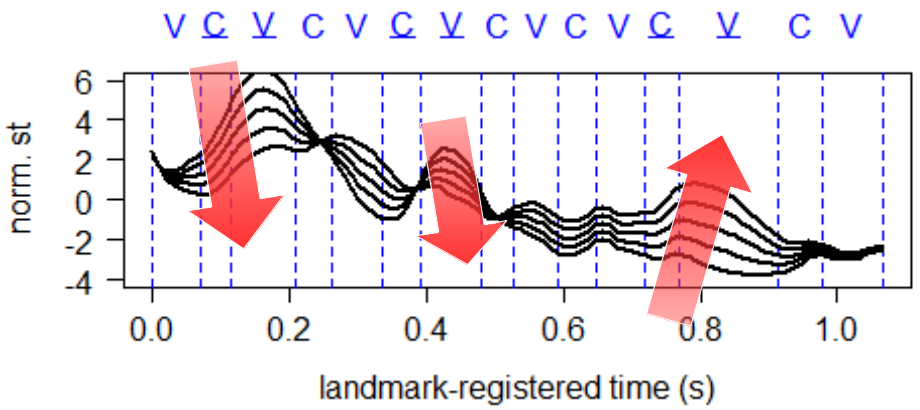

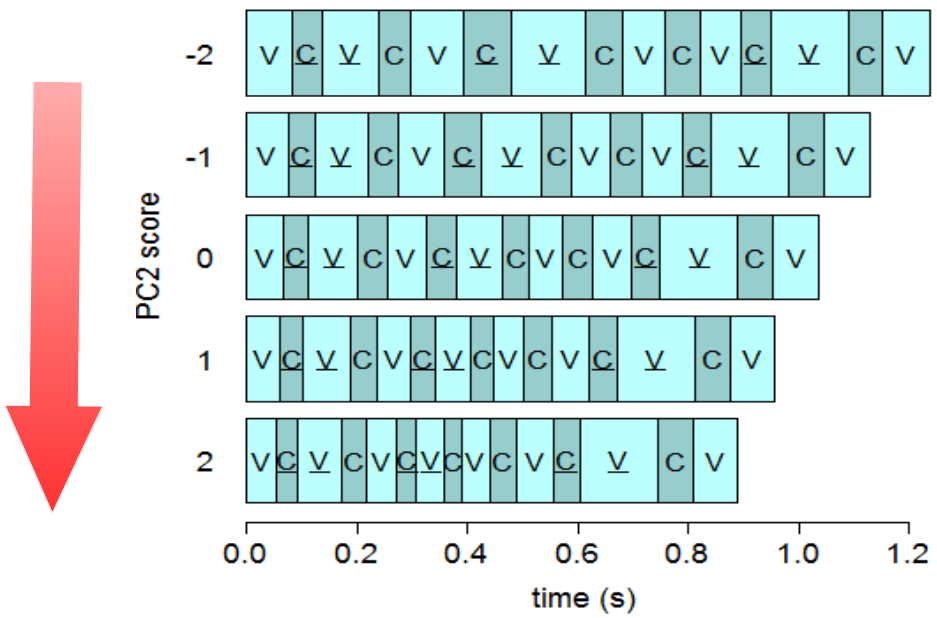

f_0

$\mu(t) + 2 * PC1(t)$
$\mu(t) + 1 * PC1(t)$
$\mu(t) + 0 * PC1(t)$
$\mu(t) - 1 * PC1(t)$
$\mu(t) - 2 * PC1(t)$

segment durations

multi-segment curve parametrisation

PC1 score

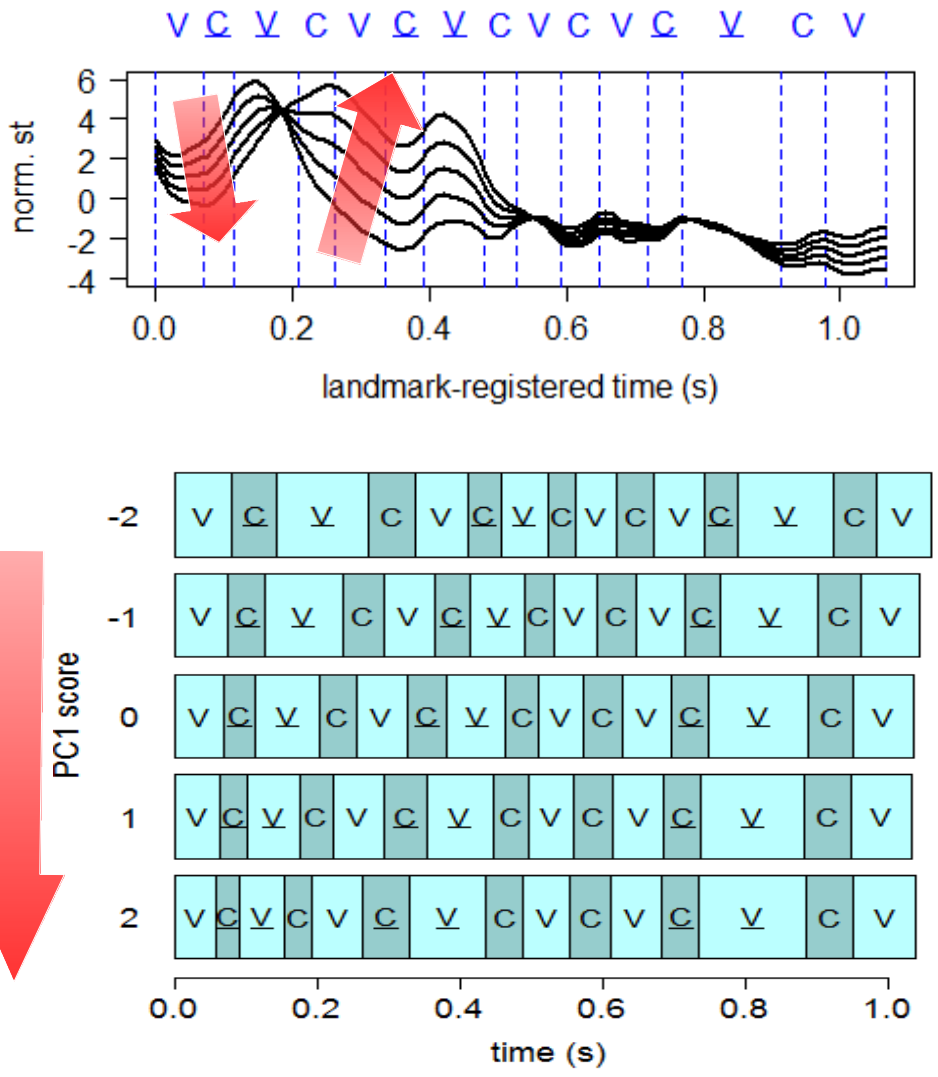

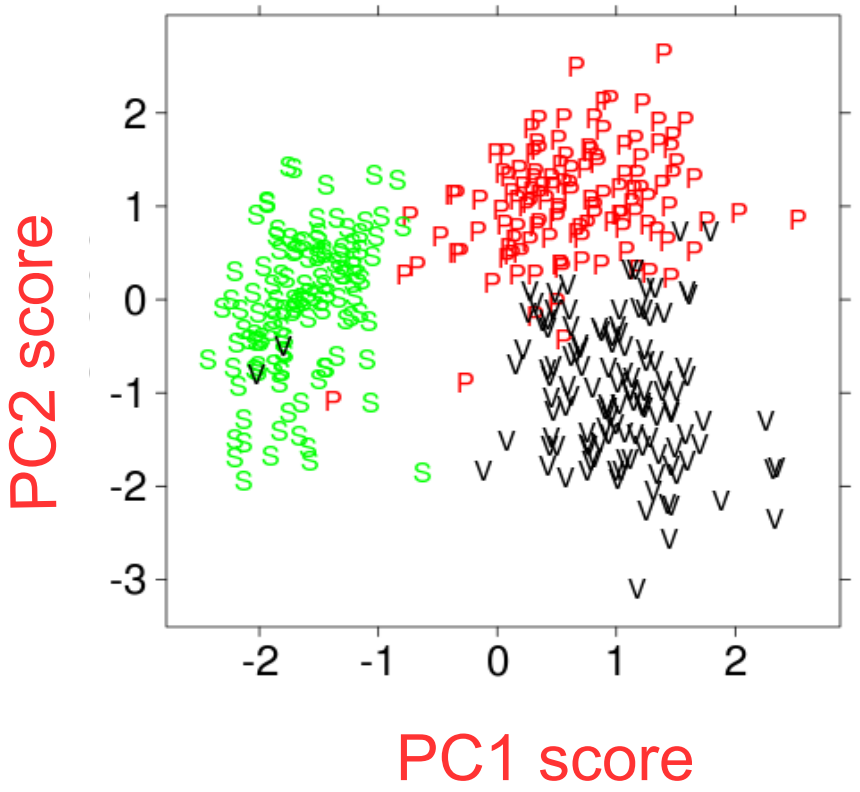

multi-segment curve parametrisation

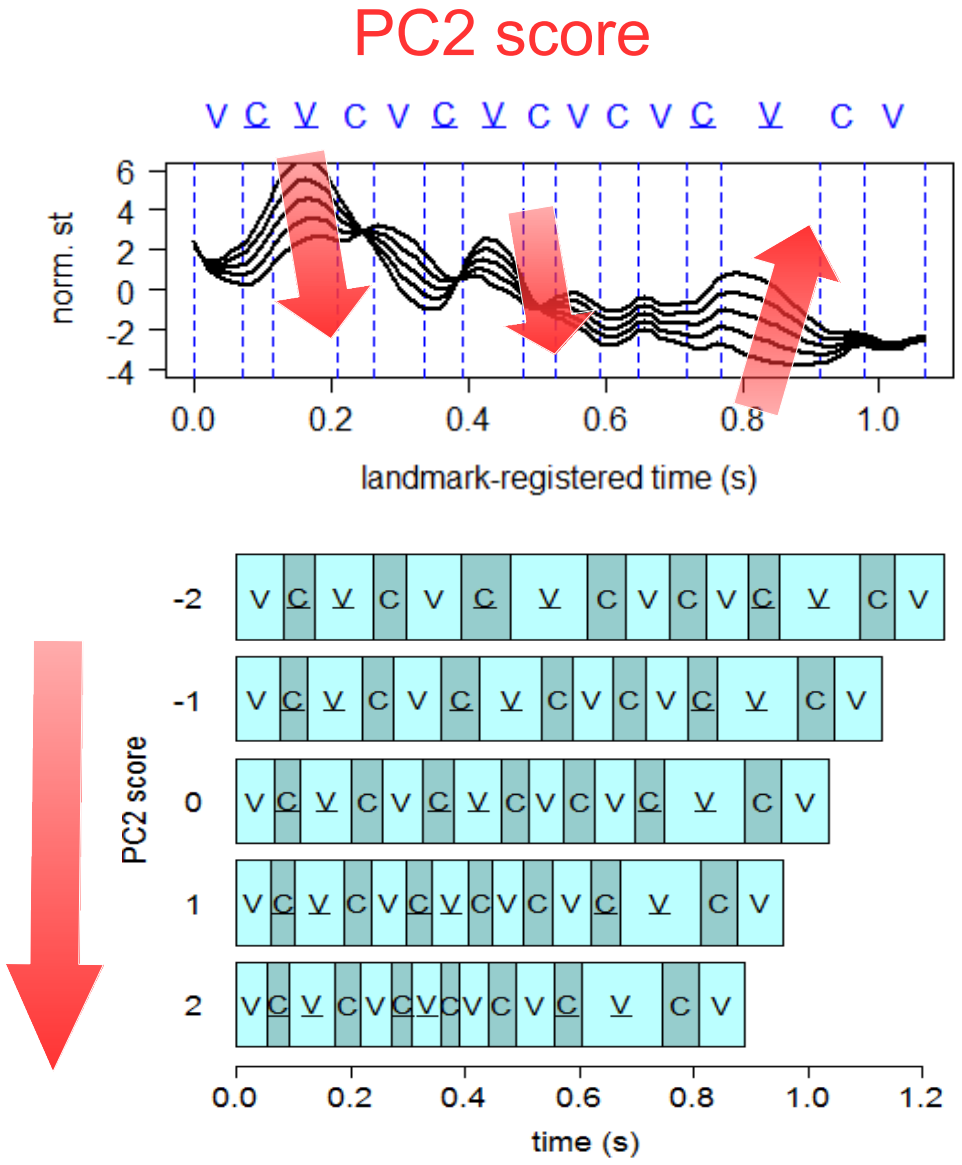

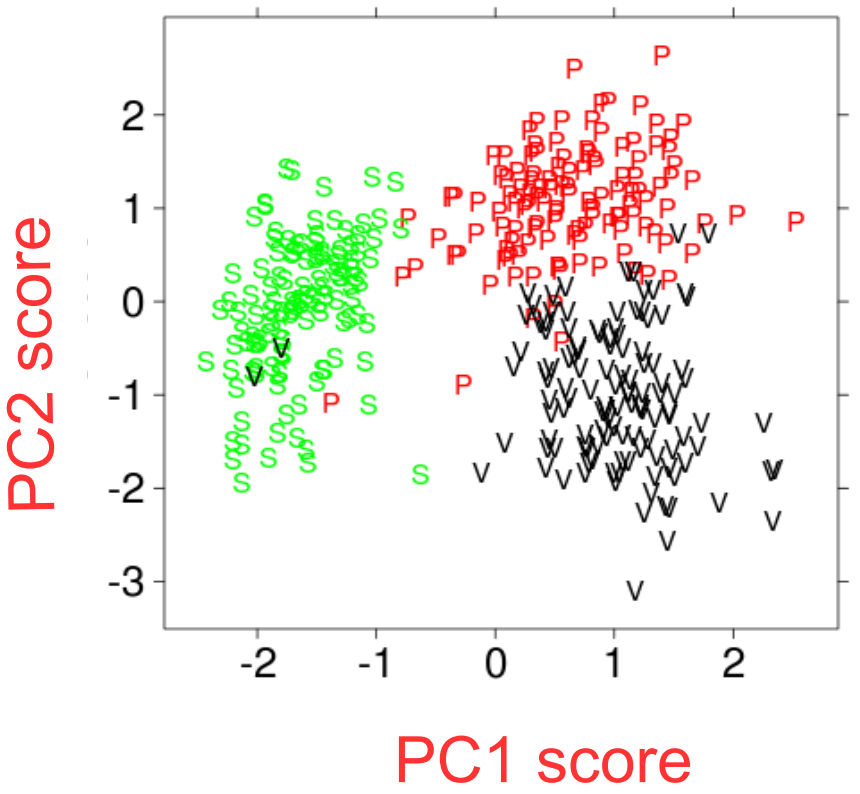

Extra slides

Take longest duration

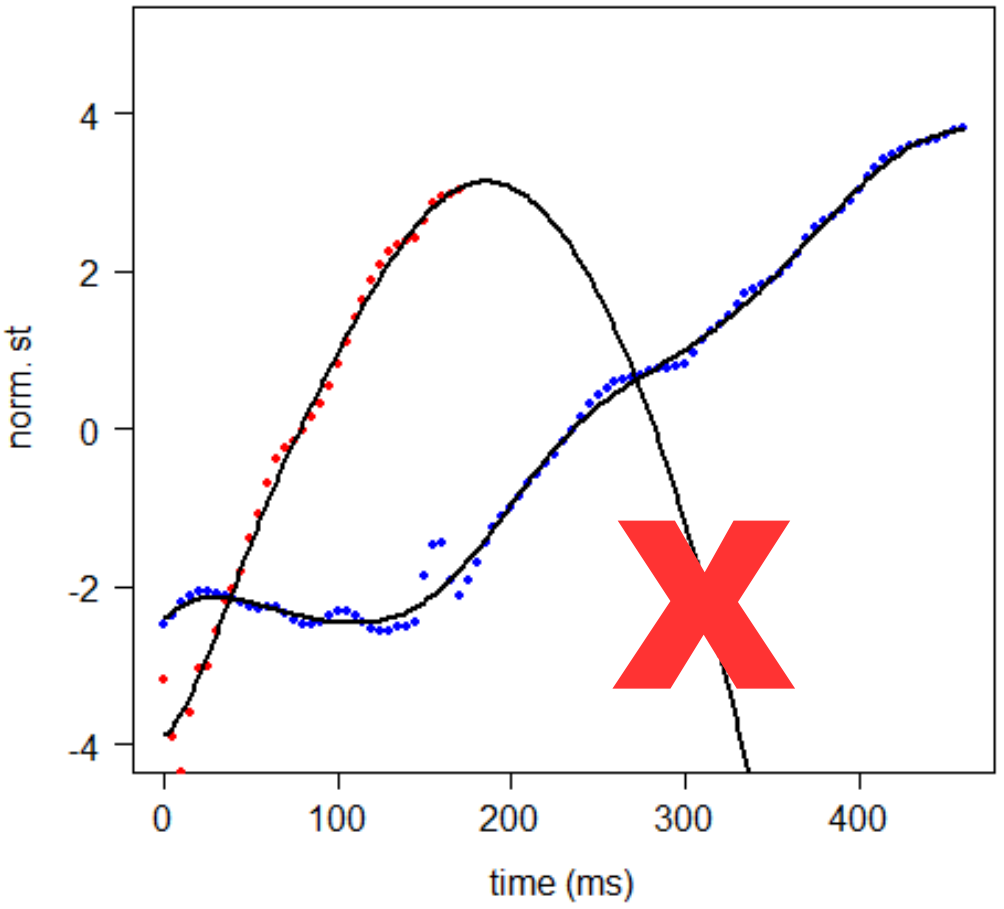

Take shortest duration

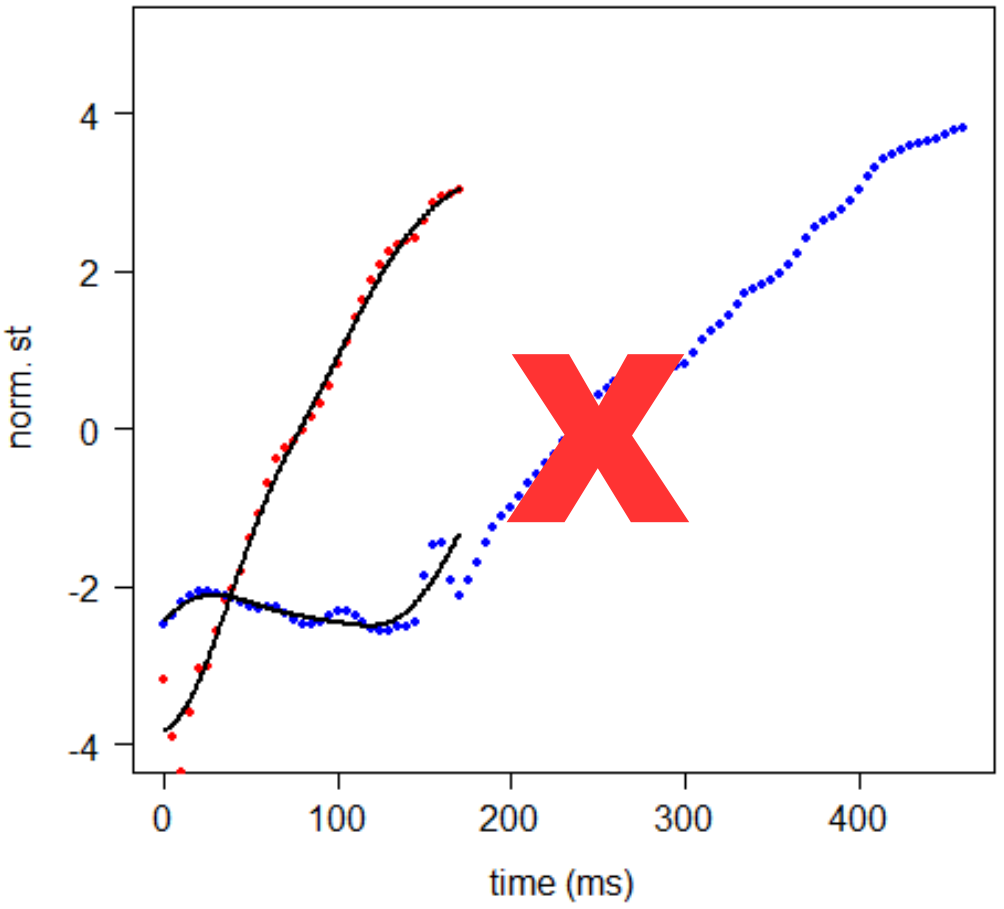

Principal Component Analysis

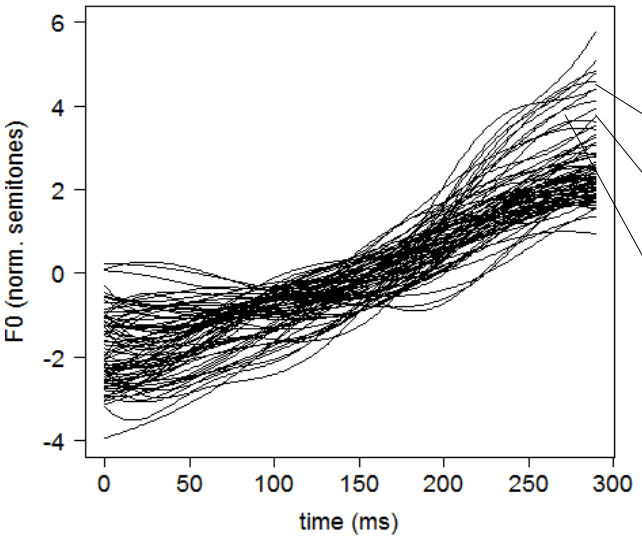

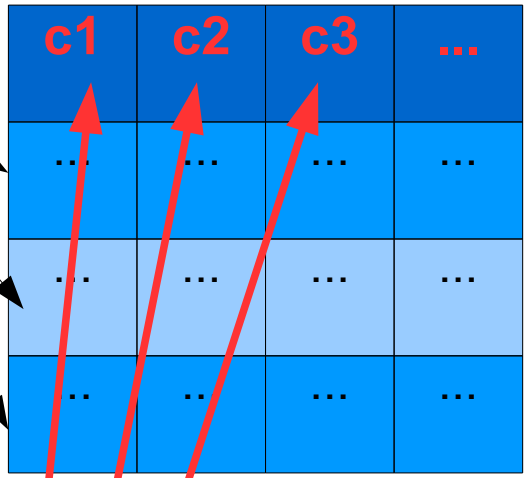

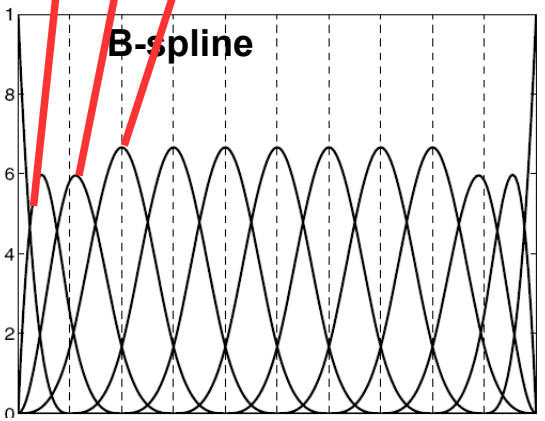

PCA

PCA limitations

- PCA does not use any explicit information related to the curve shapes or the B-splines shapes
- e.g. the sequence of coefficients c1, c2,.. reflects time adjacency of polynomial components, i.e. overlapping 'hills'

Functional PCA

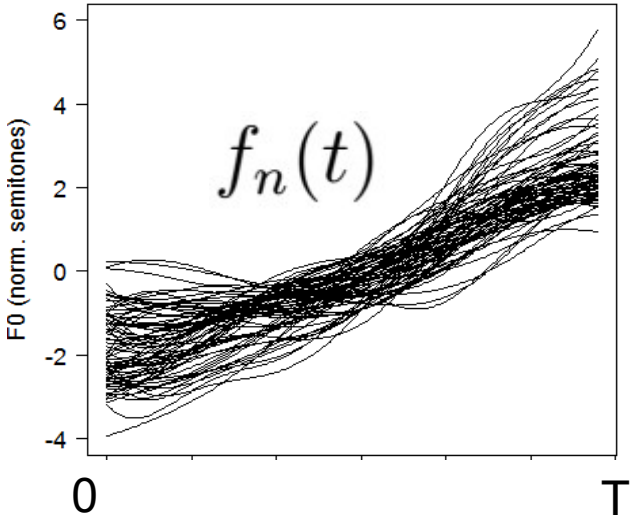

$$\max\left\{var_n\left(\int_0^T PC1(t)f_n(t)dt\right)\right\}$$

$$\text{subject to}\int_0^T PC1^2(t)=1$$

- FPCA definition uses the input curves $f_n(t)$
- FPCA is independent of the B-splines used to smooth $f_n(t)$